

**Data Structures and Algorithms  
(ES221)**

# **QUEUES (HEAPS)**

**Dr. Zubair Ahmad**

# Priority Queues



- Jobs sent to a line printer are generally placed on a queue
- In a multiuser environment, the operating system scheduler must decide which of several processes to run
- **The queue is FCFS**
  - dequeue the first job
  - run it until either it finishes or its time limit is up, and
  - enqueue if not done within the time limit
  - FCFS may not always be so efficient
    - May not be fair with the short jobs, waiting too long in the queue
- What is the solution?
- Shortest job first (SJF)
  - finish the short jobs as fast as possible
  - Shorter jobs should have preference over longer jobs running
- Priority Queues
  - Some jobs that are not short are still very important and should also have preference

# Priority Queues



- **FCFS**

- High Priority jobs may be delayed
- A long job may force several smaller jobs to wait for unreasonably long period of time

Pages	100	4	2	2	2	1	111
Wait time	0	100	104	106	108	110	528
Printing Time	100	104	106	108	110	111	539

Average waiting time =  $528 / 6 = 88$  units

Average printing time =  $539 / 6 = 89.8$  units

- **SJF (Shortest Job First)**

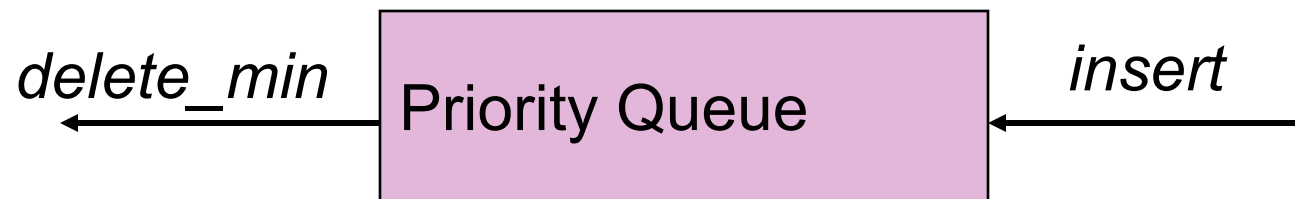
Pages	1	2	2	2	4	100	111
Wait time	0	1	3	5	7	11	27
Printing Time	1	3	5	7	11	111	138

Average waiting time =  $27 / 6 = 4.5$  units

Average printing time =  $138 / 6 = 23$  units

# Priority Queues Implementation

- Two usual operations
  - *Insert*. Insert a new job in the queue
  - *delete\_min*. Remove the shortest/highest priority job from the queue



# Priority Queues Implementation

- Simple linked list implementation
  - ***Insert:*** Insert at the front/end of the queue
    - $O(1)$  running time
  - ***Delete\_min:*** requires to traverse the whole list
    - $O(n)$  running time
- Sorted linked list implementation
  - ***Insert:*** Insert at its proper location requires traversal
    - $O(n)$  running time
  - ***Delete\_min:*** delete the first element which is the min.
    - $O(1)$  running time

# Priority Queues Implementation

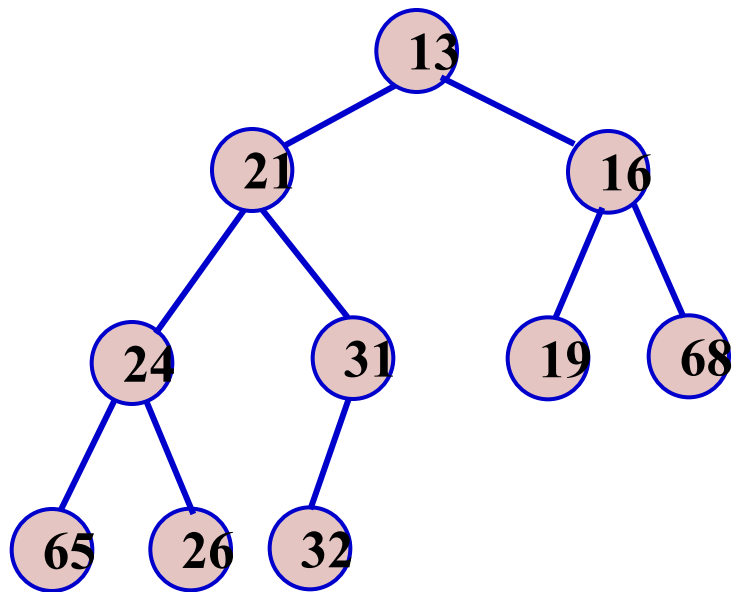
- Binary search tree implementation
  - **Insert:** Insert at it proper location requires tree traversal
    - $O(\log n)$  running time
  - **Delete\_min:** delete the deepest child from the leftmost subtree
    - $O(\log n)$  running time
- Insertions are random
- deletions are not (why?)
  - Repeated removal of the minimum will hurt the balance of the tree (how?)
  - Worst case scenario
    - The left subtree is depleted (all its nodes deleted)
    - the right subtree would have at most twice as many elements as it should.
    - This adds only a small constant to its expected depth
- Using Binary Search is over skill
  - We only need two operations

# Heaps

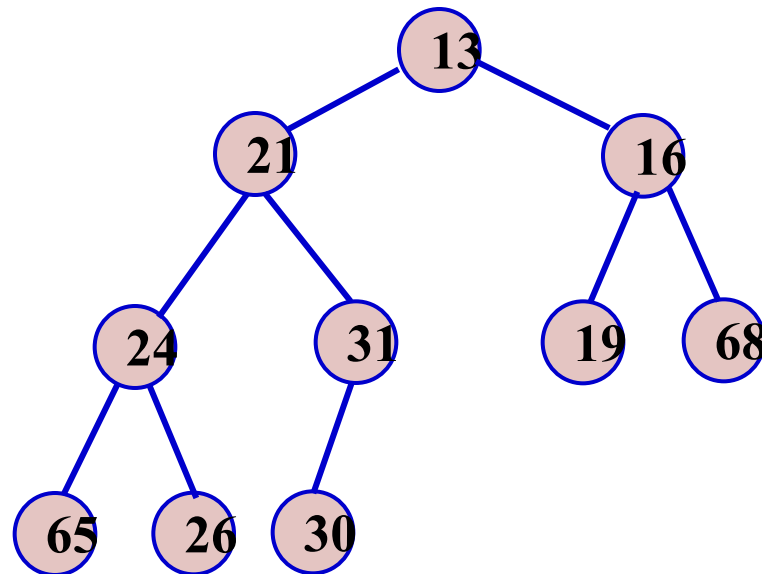
- *binary heap ~ heap*
- *A binary tree which is completely filled*
- *The last level may not be completely filled*
  - It is filled from left to right
- It is easy to show that a complete binary tree of height  $h$  has between  $2^h$  and  $2^{h+1} - 1$  nodes
- The height of a complete binary tree is  $\lceil \log n \rceil$  which is clearly  $O(\log n)$ .

# Heaps

Value of an element at a node is less than or equal to that of its descendants. (Heap property)



**A heap**



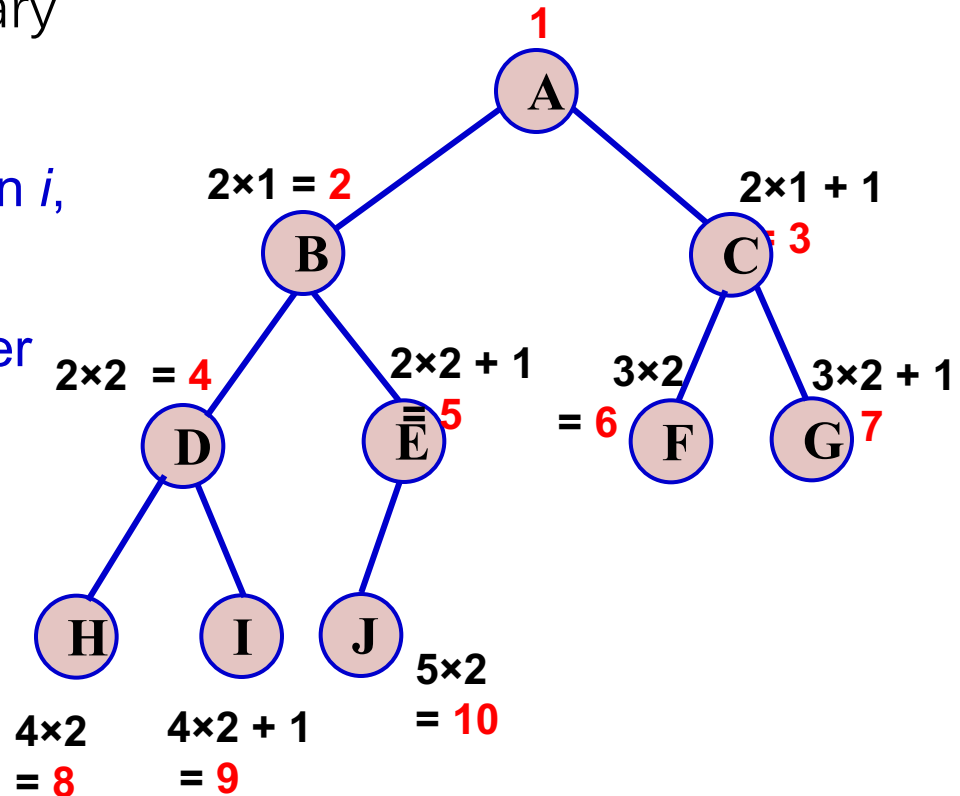
**Not a heap**



# Heaps: Array implementation

- An important observation is that because a complete binary tree is so regular, it can be represented in an array and no pointers are necessary

- Start from index 1
- For any element in array position  $i$ ,
  - the left child is in position  $2i$ ,
  - the right child is in the cell after the left child ( $2i + 1$ ), and
  - the parent is in position  $i/2$
- Any problem?
  - an estimate of the maximum heap size is required in advance



	A	B	C	D	E	F	G	H	I	J			
0	1	2	3	4	5	6	7	8	9	10	11	12	13

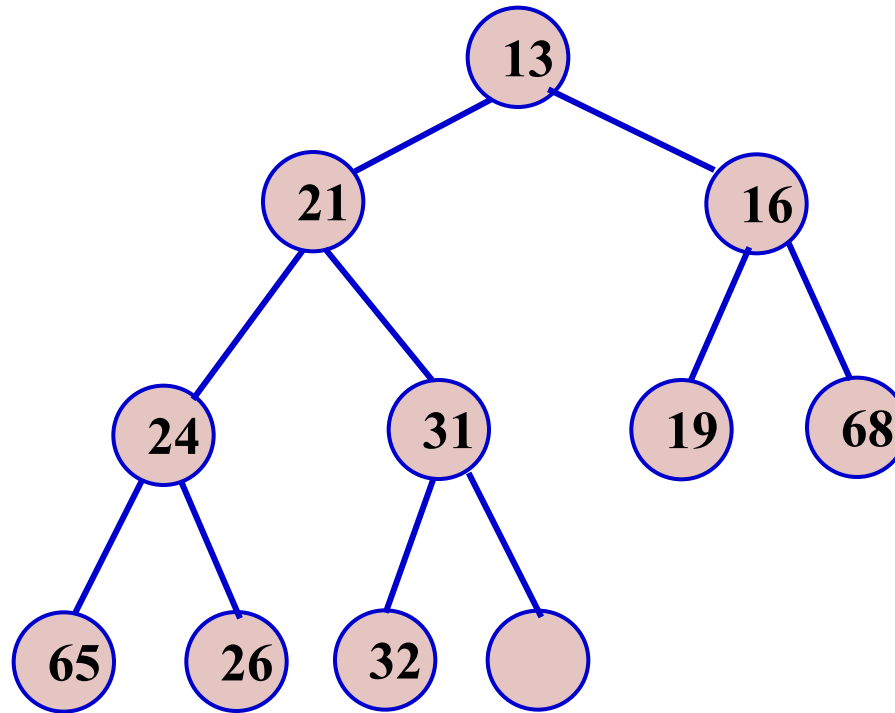
# Heaps - Operations

- Insert
- Delete

# Heaps – insert operation

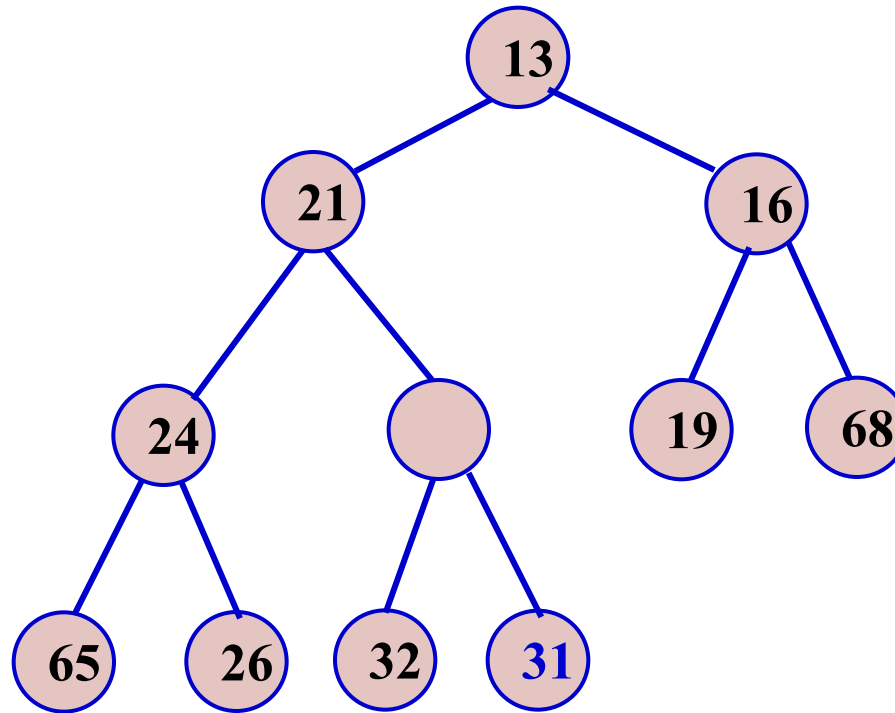
- *Insert  $x$  into the heap*
  - *Create a hole in the next available location*
    - *since otherwise the tree will not be complete*
  - *If  $x$  can be placed in the hole without violating heap order, then*
    - *we do so and are done*
  - *Else*
    - *bubble the hole up*
      - *by sliding the element that is in the hole's parent node into the hole*
  - *Continue this process until  $x$  can be placed in the hole*
  - *This general strategy is known as a *percolate up*;*
    - *the new element is percolated up the heap until the correct location is found*

# Heaps – insert 14



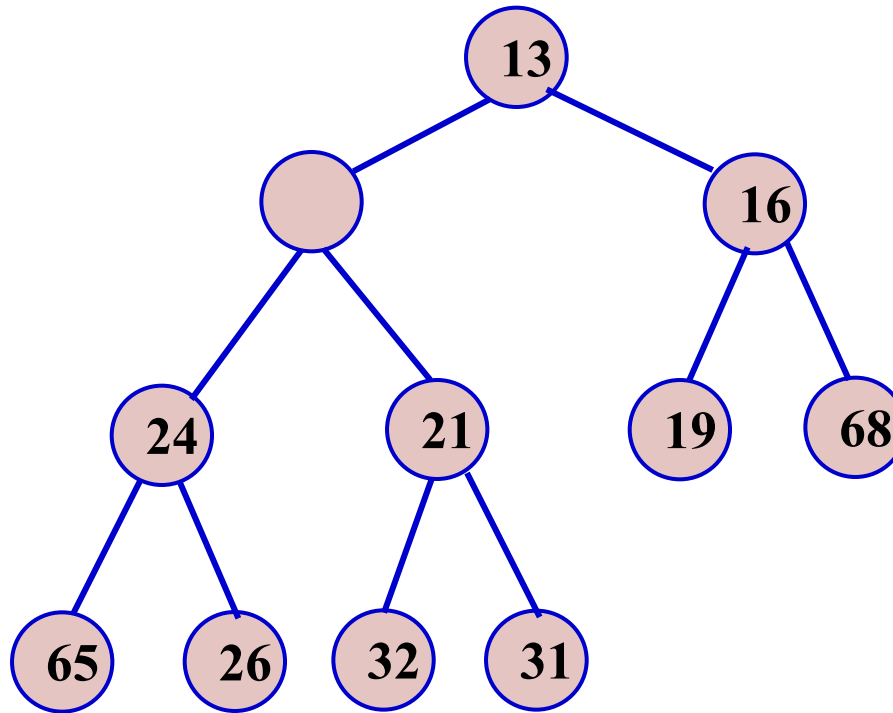
	13	21	16	24	31	19	68	65	26	32			
0	1	2	3	4	5	6	7	8	9	10	11	12	13

# Heaps – insert 14



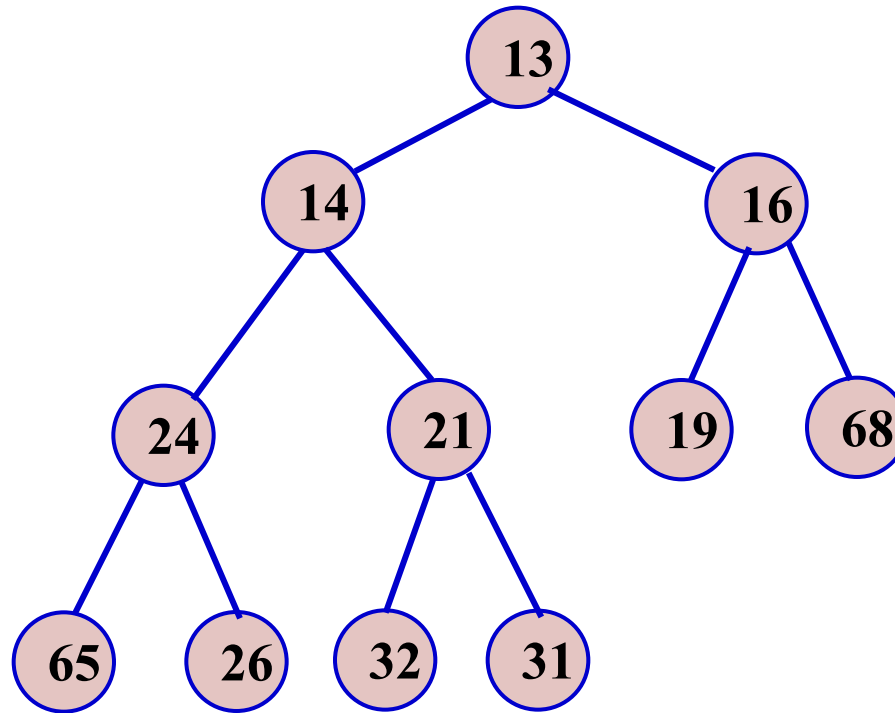
	13	21	16	24		19	68	65	26	32	31		
0	1	2	3	4	5	6	7	8	9	10	11	12	13

# Heaps – insert 14



	13		16	24	21	19	68	65	26	32	31		
0	1	2	3	4	5	6	7	8	9	10	11	12	13

# Heaps – insert 14



	13	14	16	24	21	19	68	65	26	32	31		
0	1	2	3	4	5	6	7	8	9	10	11	12	13

# Heaps – insert Operation

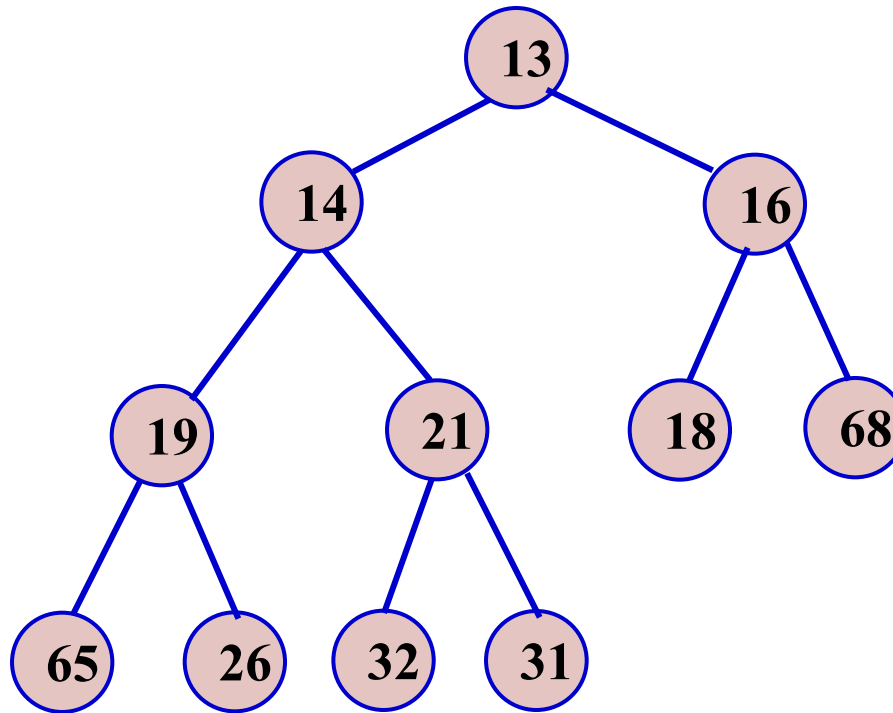
- We could have implemented the percolation in the insert routine by performing repeated swaps until the correct order was established
  - but a swap requires three assignment statements.
- If an element is percolated up  $d$  levels, the number of assignments performed by the swaps would be  $3d$



# Heaps – delete\_min Operation

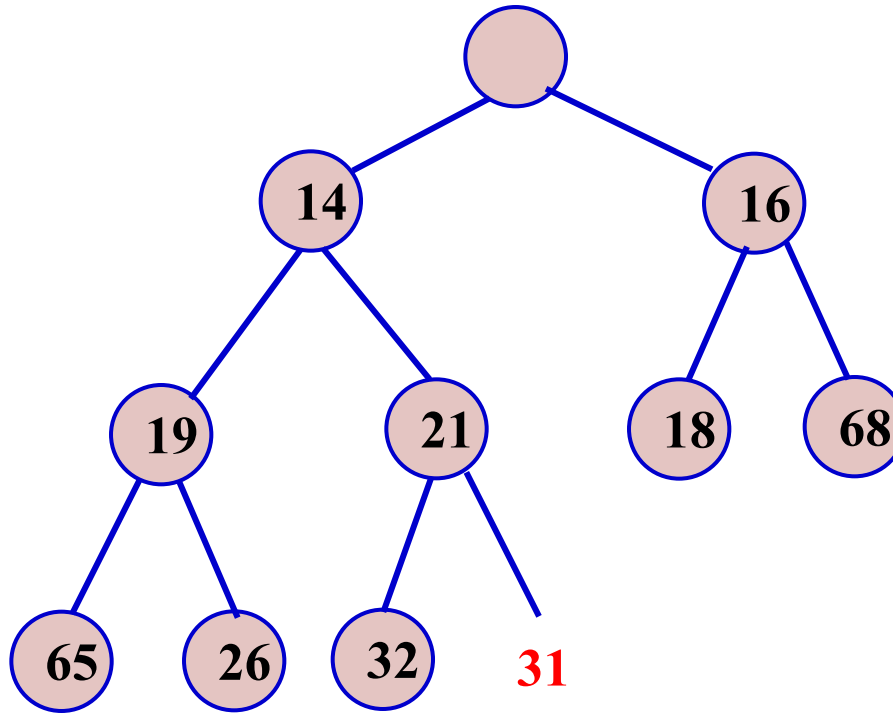
- *Analogous to* insertions.
- When the minimum is removed, a hole is created at the root.
- Since the heap now becomes one smaller, it follows that the last element *x* in the heap must move somewhere in the heap
- If *x* can be placed in the hole (which is unlikely) then
  - we are done
- Else
  - push the hole down one level
    - by sliding the smaller of the hole's children into the hole
- Repeat this step until *x* can be placed in the hole
- **The result is:** place *x* in its correct spot along a path from the root containing *minimum* children
  - This general strategy is known as a *percolate down*;

# Heaps – delete\_min : 13



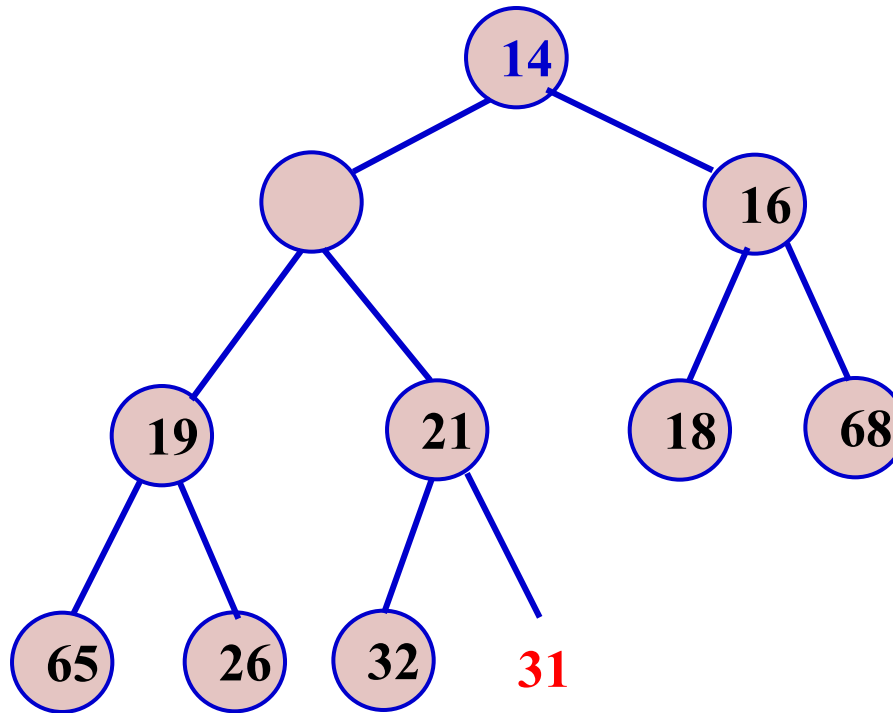
	13	14	16	19	21	18	68	65	26	32	31		
0	1	2	3	4	5	6	7	8	9	10	11	12	13

# Heaps – delete\_min : 13



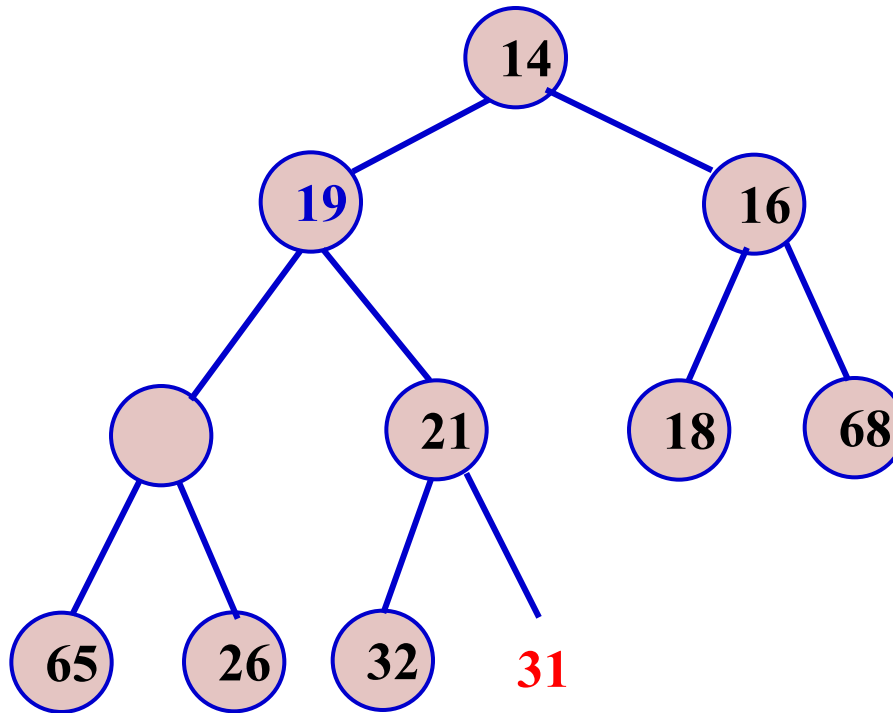
		14	16	19	21	18	68	65	26	32	31		
0	1	2	3	4	5	6	7	8	9	10	11	12	13

# Heaps – delete\_min : 13



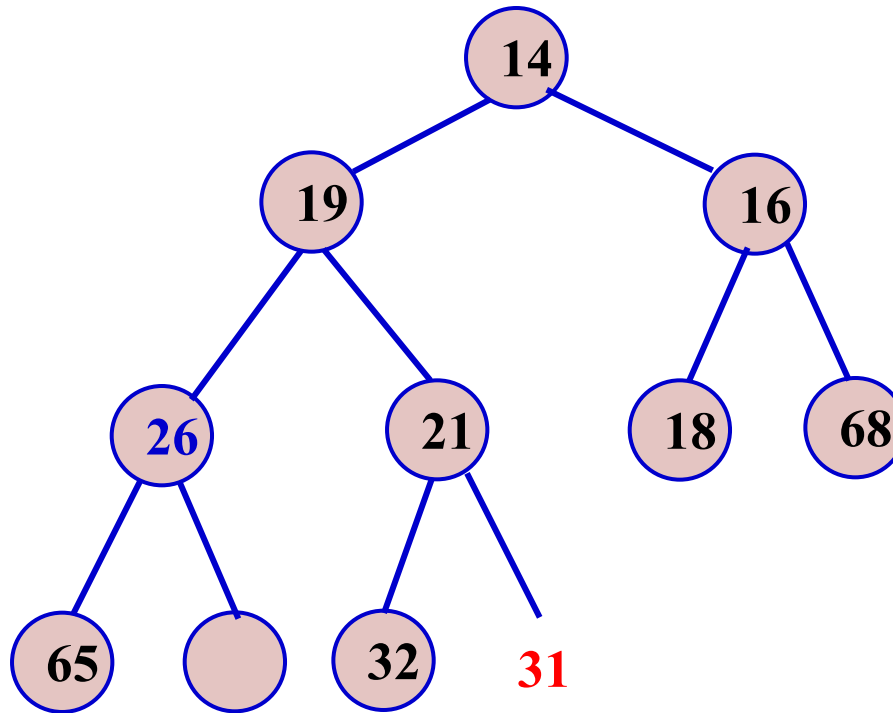
	14		16	19	21	18	68	65	26	32	31		
0	1	2	3	4	5	6	7	8	9	10	11	12	13

# Heaps – delete\_min : 13



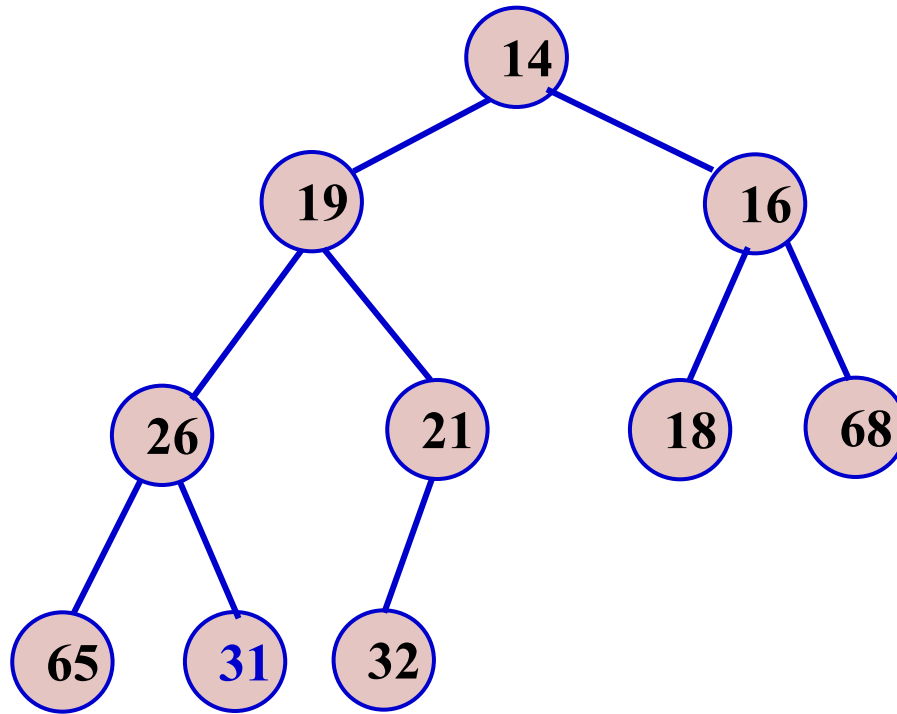
	14	19	16		21	18	68	65	26	32	31		
0	1	2	3	4	5	6	7	8	9	10	11	12	13

# Heaps – delete\_min : 13



	14	19	16	26	21	18	68	65		32	31		
0	1	2	3	4	5	6	7	8	9	10	11	12	13

# Heaps – delete\_min : 13



	14	19	16	26	21	18	68	65	31	32			
0	1	2	3	4	5	6	7	8	9	10	11	12	13

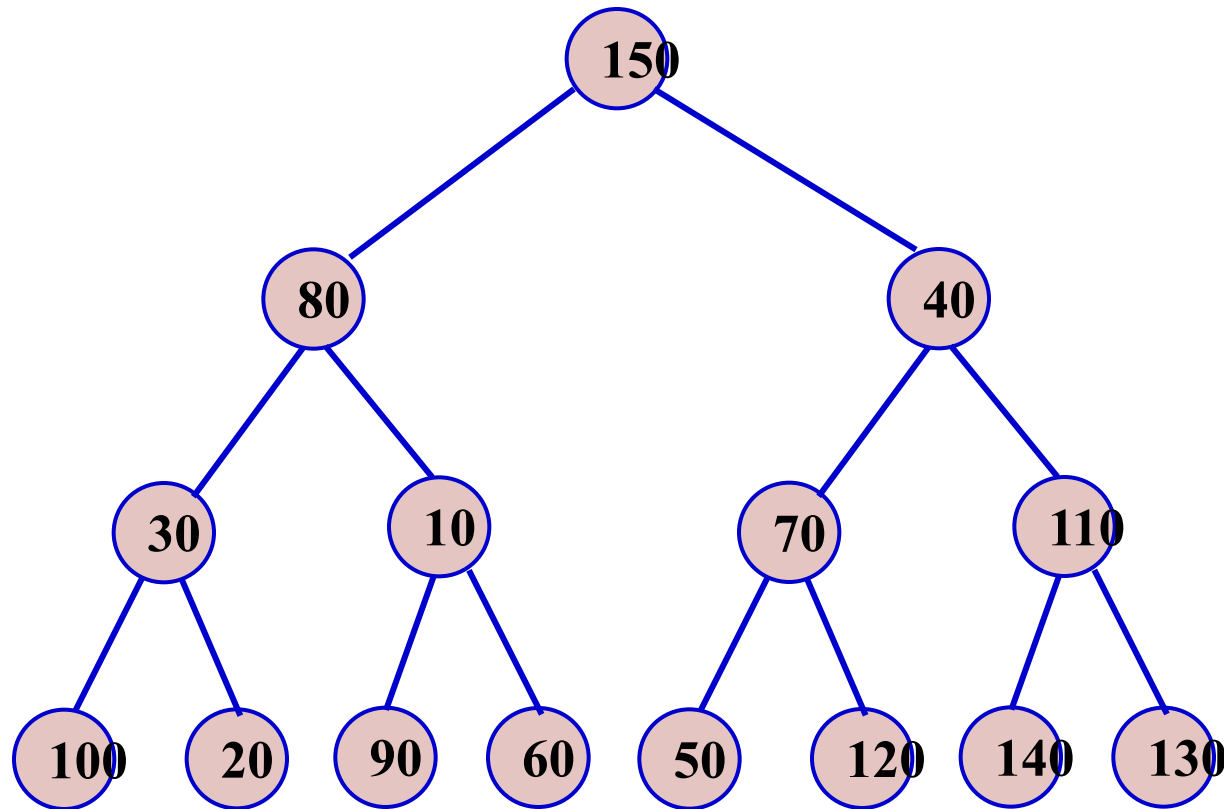
# Heap Operations: Build Heap

- The general algorithm is to place the  $n$  keys into the tree in any order, maintaining the structure property.
- Then, if *percolate\_down*( $i$ ) percolates down from node  $i$ , perform the following algorithm to create a heap-ordered tree

```
for( $i = n / 2; i > 0; i--$  )  
    percolate_down(  $i$ );
```

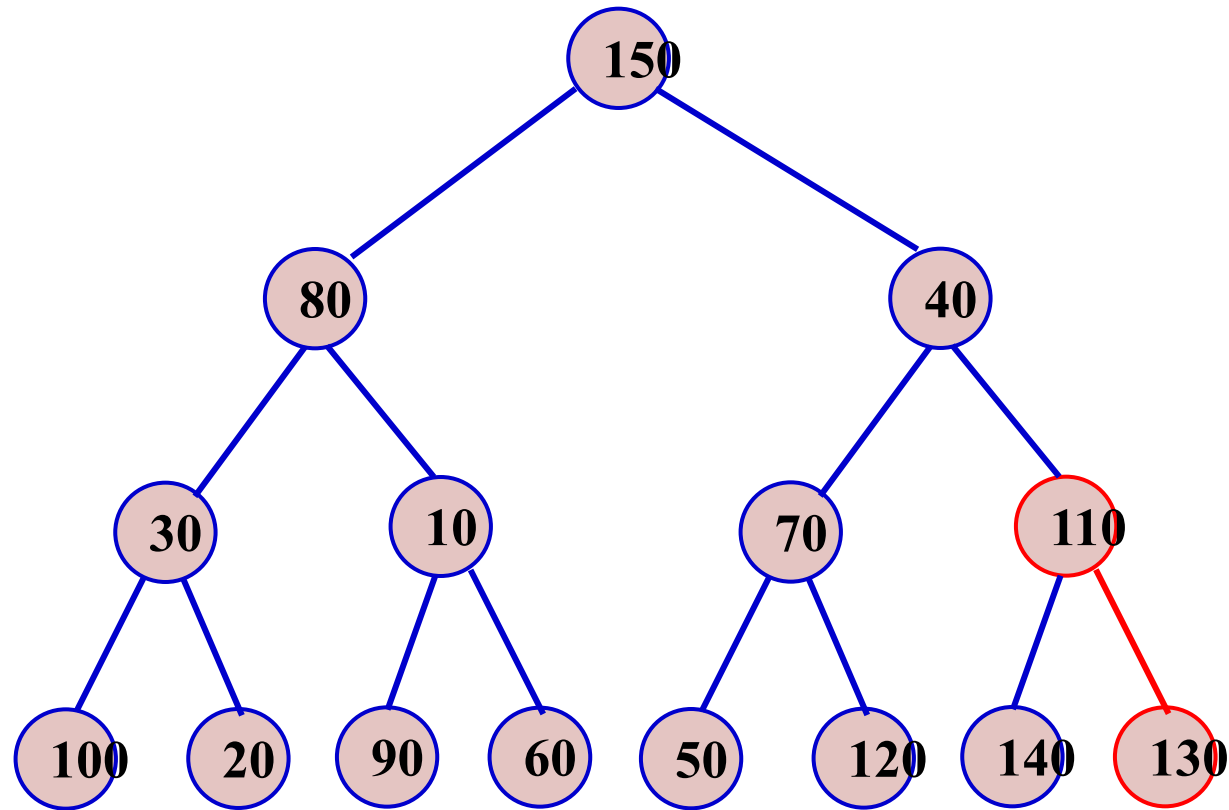


# Build Heap



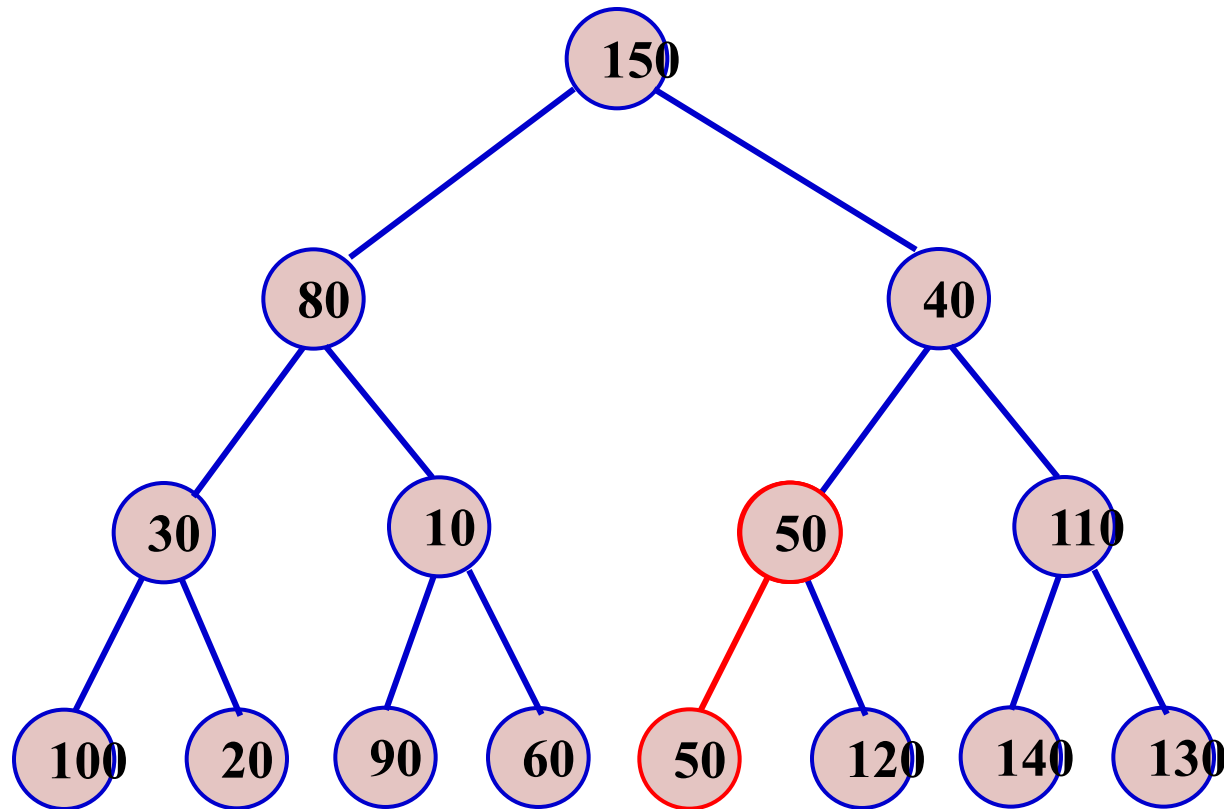
Initial (unordered) Heap

# Build Heap



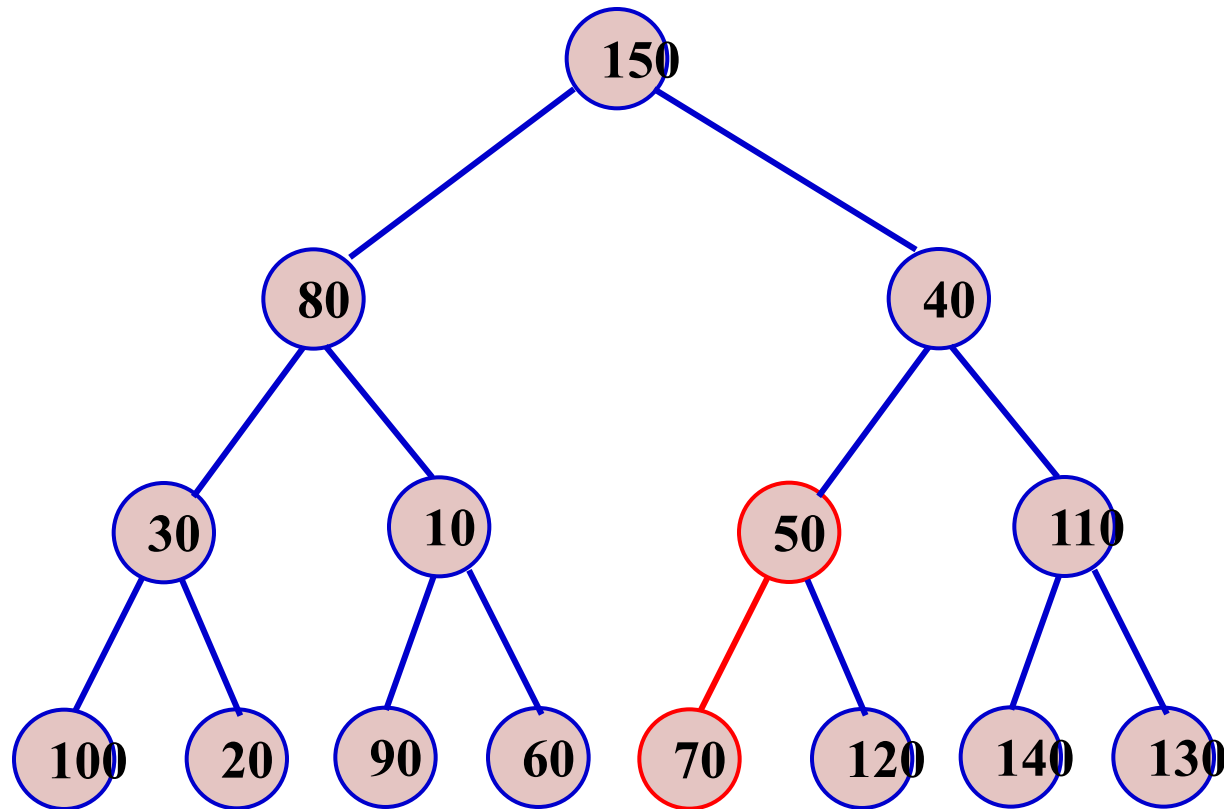
Heap after percolate\_down(7)

# Build Heap



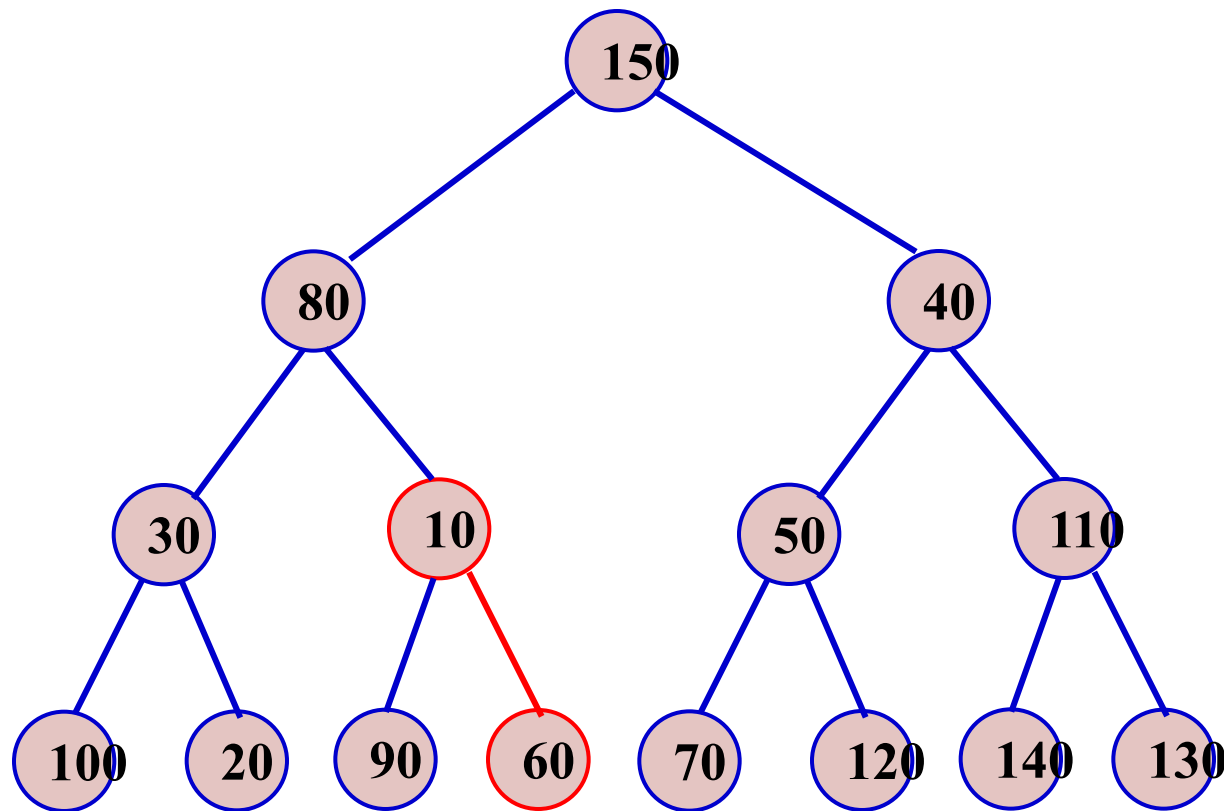
Heap after percolate\_down(6)

# Build Heap



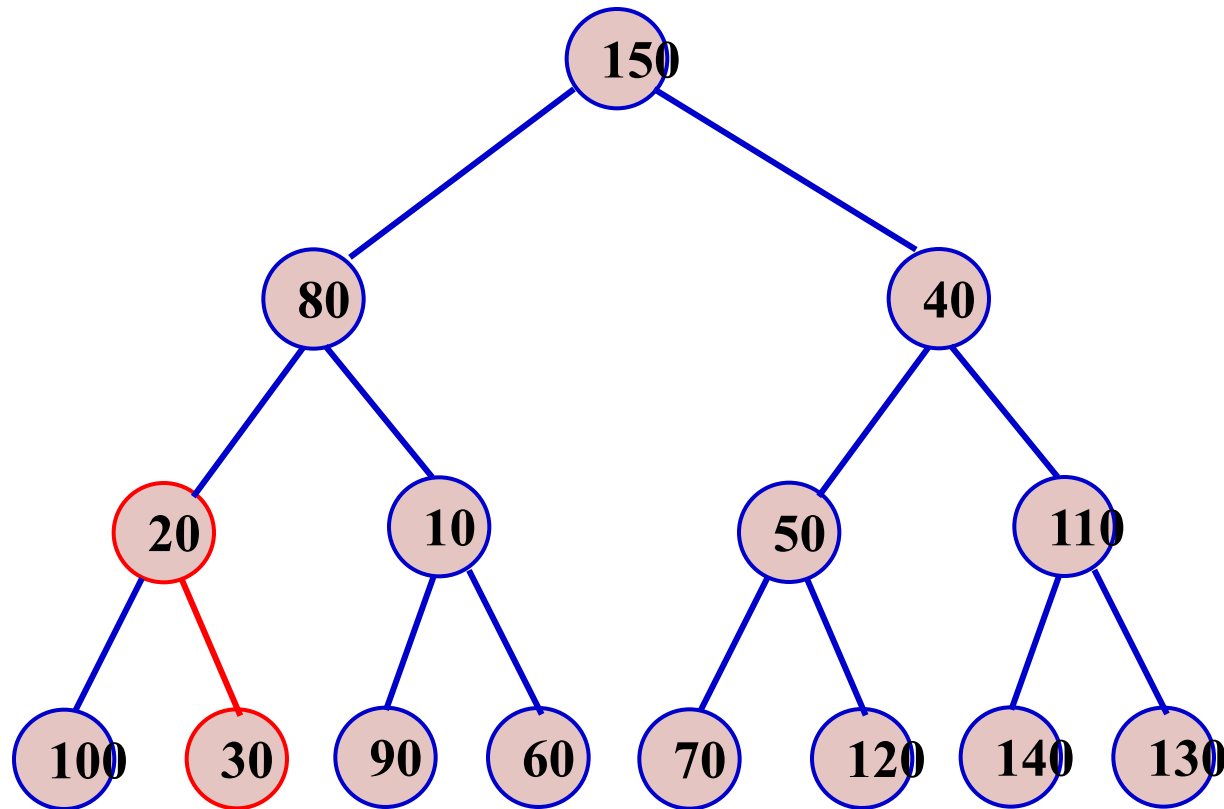
Heap after percolate\_down(6)

# Build Heap



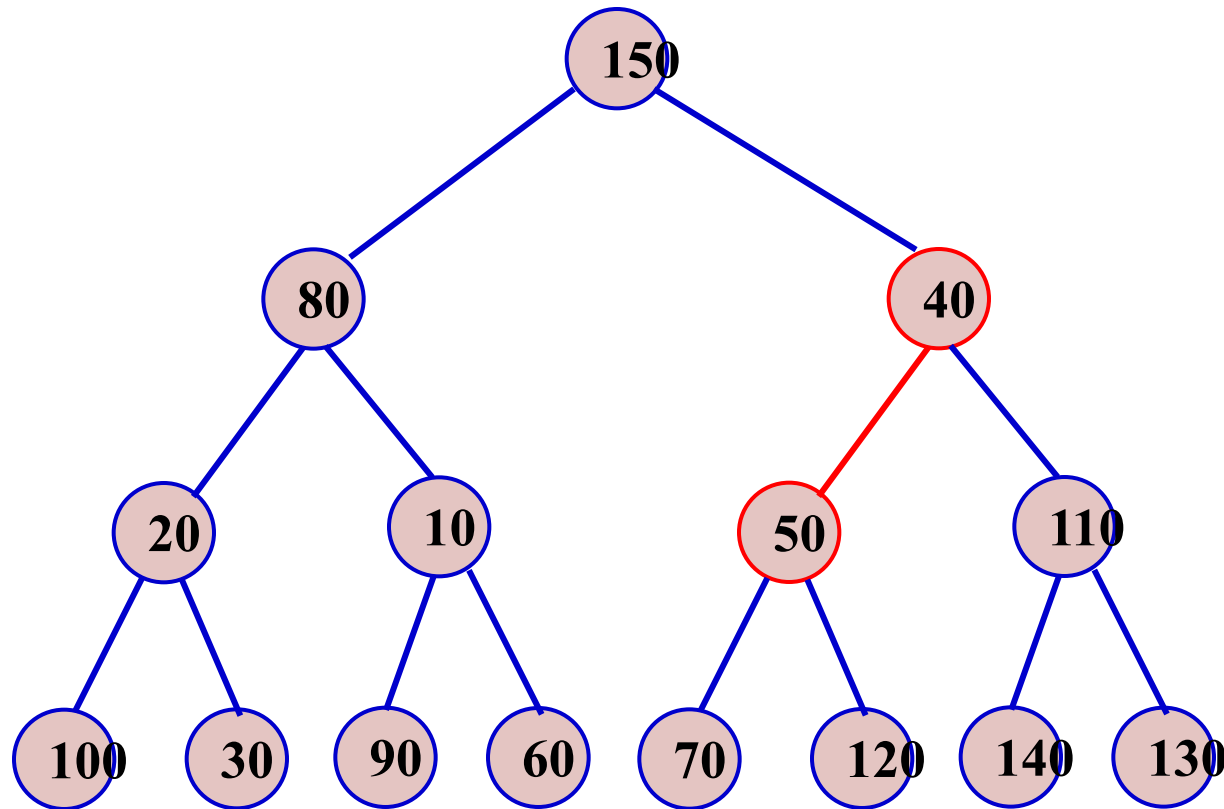
Heap after percolate\_down(5)

# Build Heap



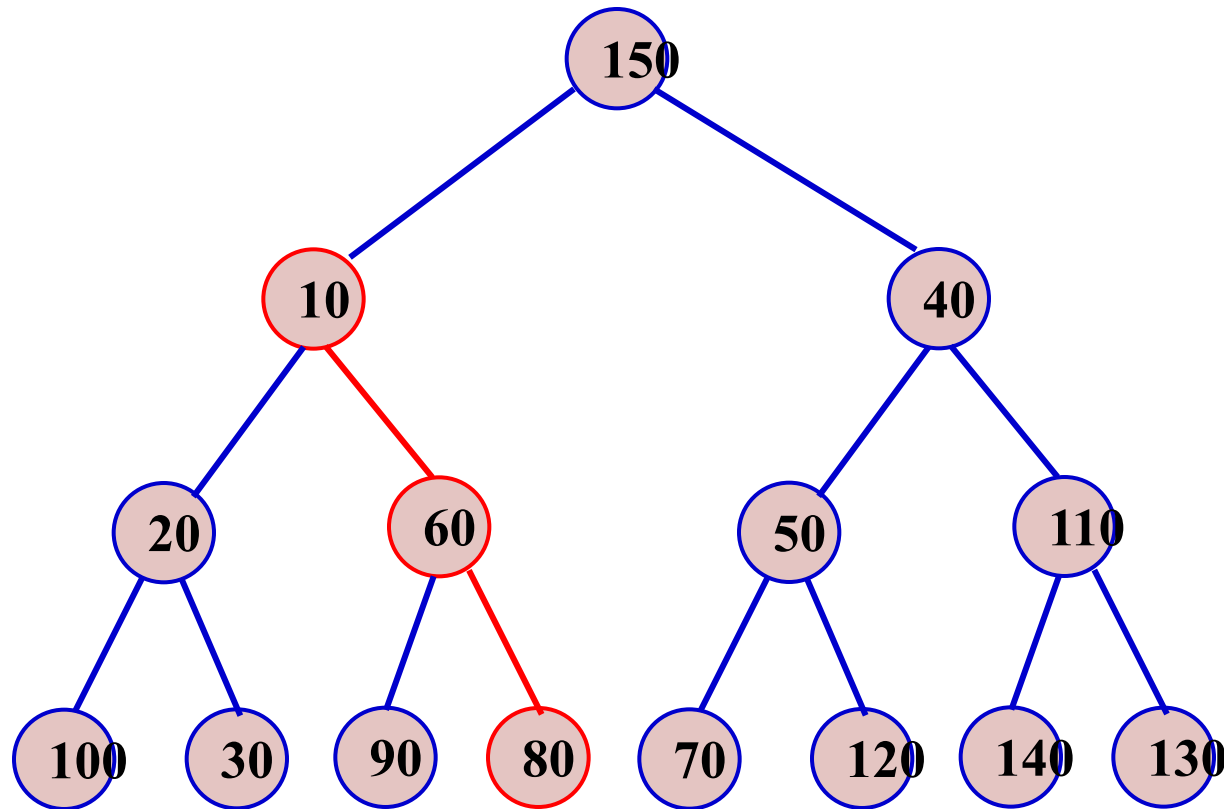
Heap after percolate\_down(4)

# Build Heap



Heap after percolate\_down(3)

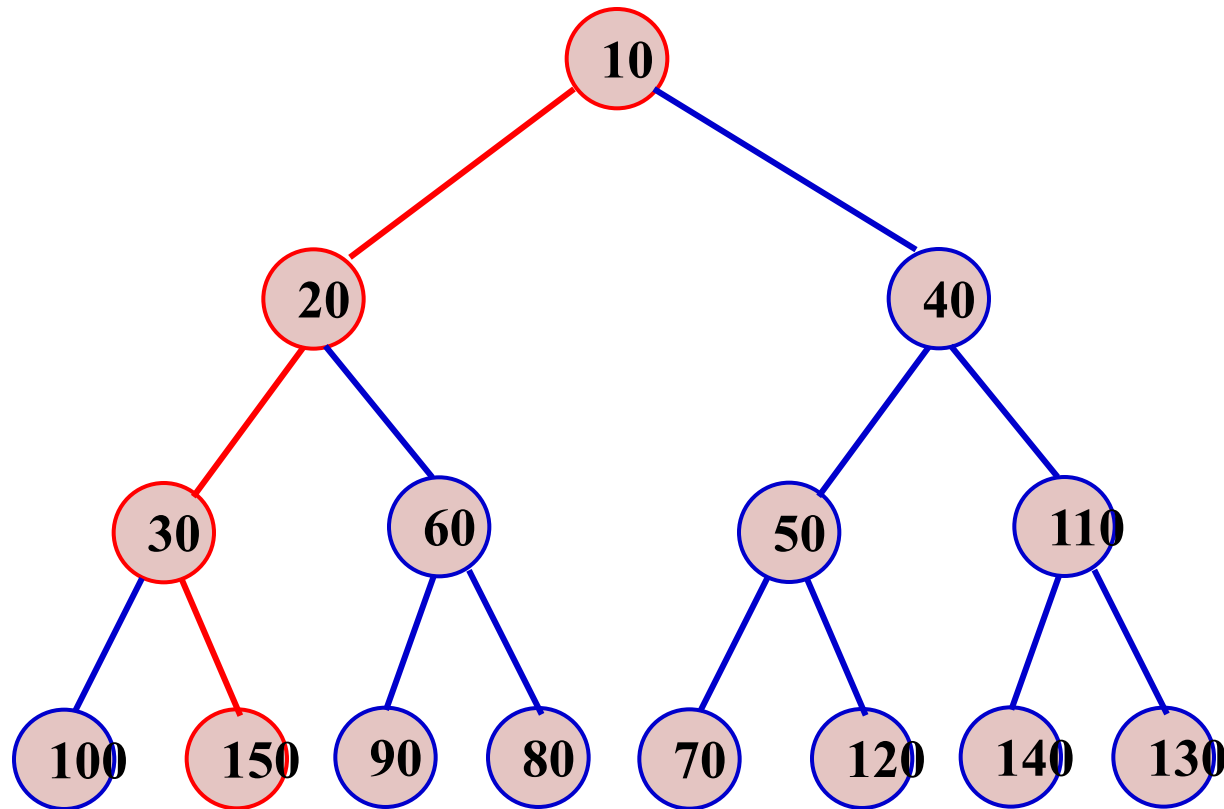
# Build Heap



Heap after percolate\_down(2)

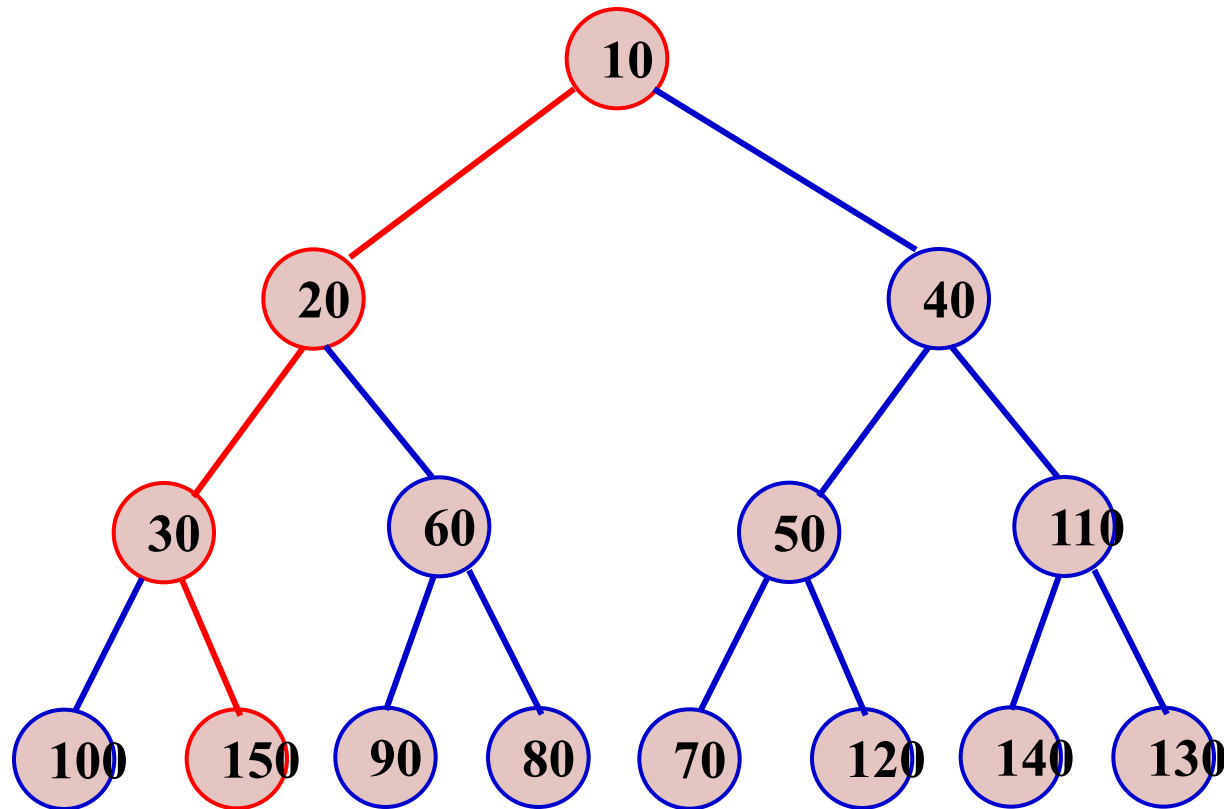


# Build Heap



Heap after percolate\_down(1)

# Build Heap



Heap after percolate\_down(1)

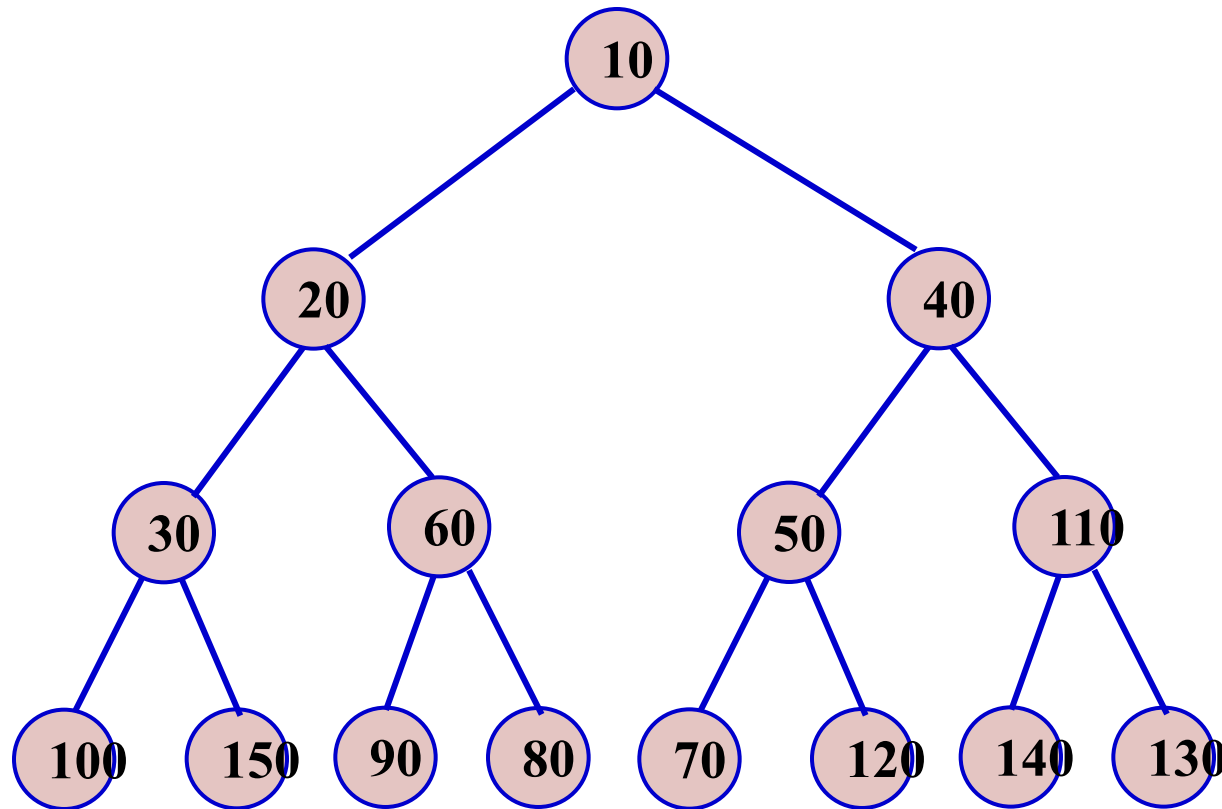
# Heap Operations: Decrease Key

- The *decrease\_key*( $x$ ,  $del$ ,  $H$ ) operation lowers the value of the key at position  $x$  by a positive amount *del*.
- Since this might violate the heap order, it must be fixed by a *percolate up*.
- This operation could be useful to system administrators:
  - they can make their programs run with highest priority

# Heap Operations: Increase Key

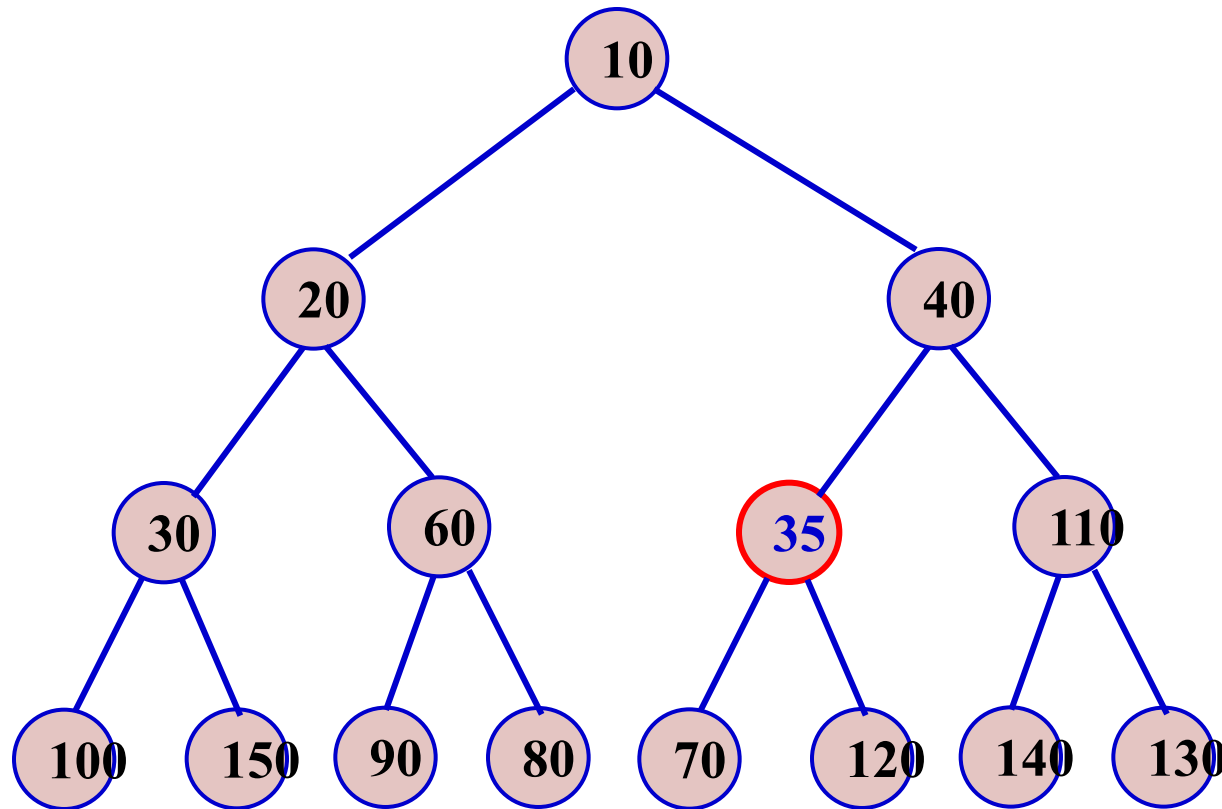
- The ***increase\_key**( $x$ ,  $del$ ,  $H$ )* operation increases the value of the key at position  $x$  by a positive amount ***del***.
- This is done with a ***percolate down***.
- Many schedulers automatically drop the priority of a process that is consuming excessive CPU time.

# Heap Operations: Decrease Key



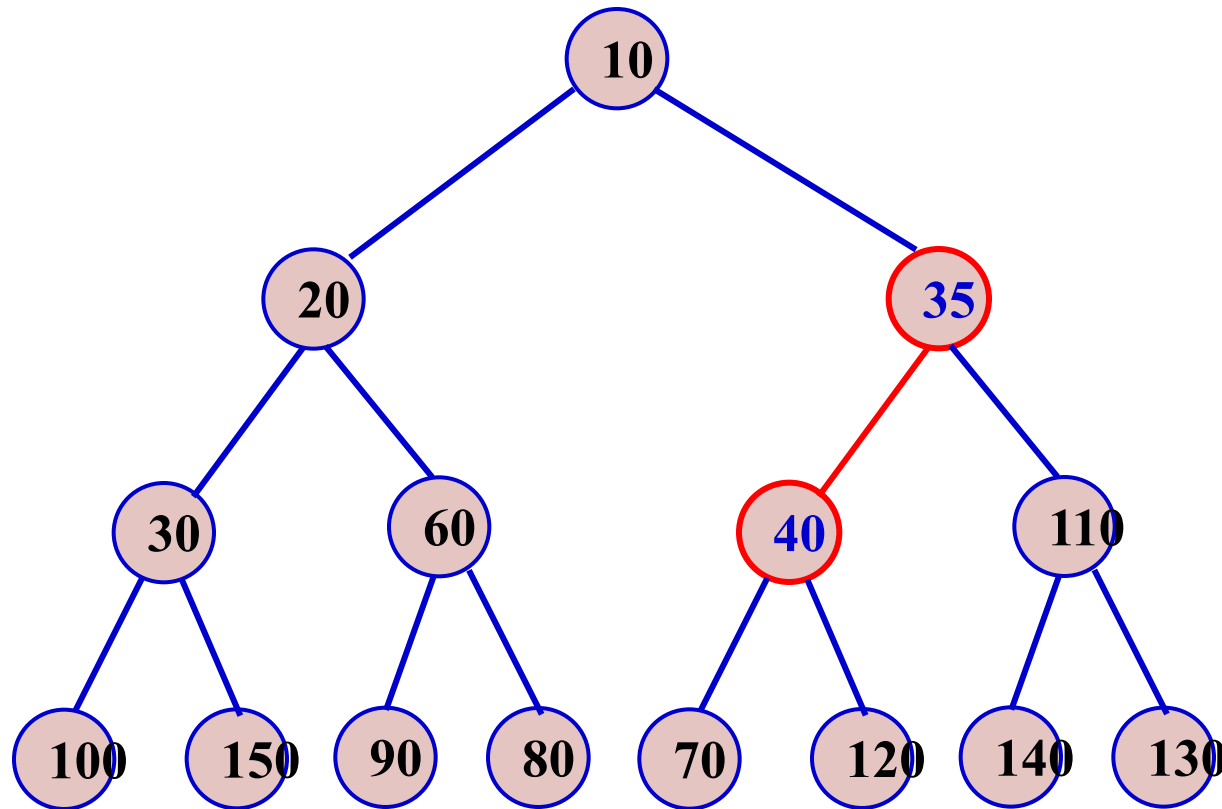
***decrease\_key(6, 15, H)***

# Heap Operations: Decrease Key



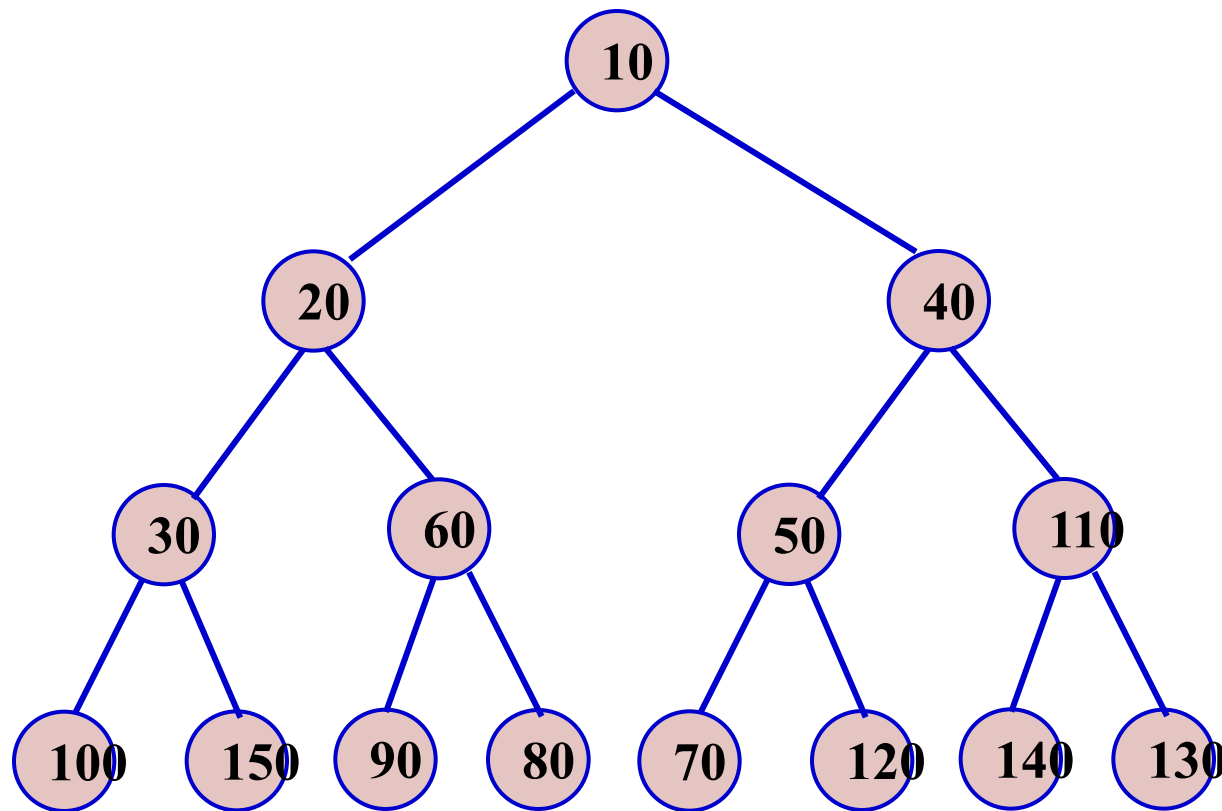
Heap decrease key(6, 15, up)(6)

# Heap Operations: Decrease Key



**Heap after percolate\_up(6)**

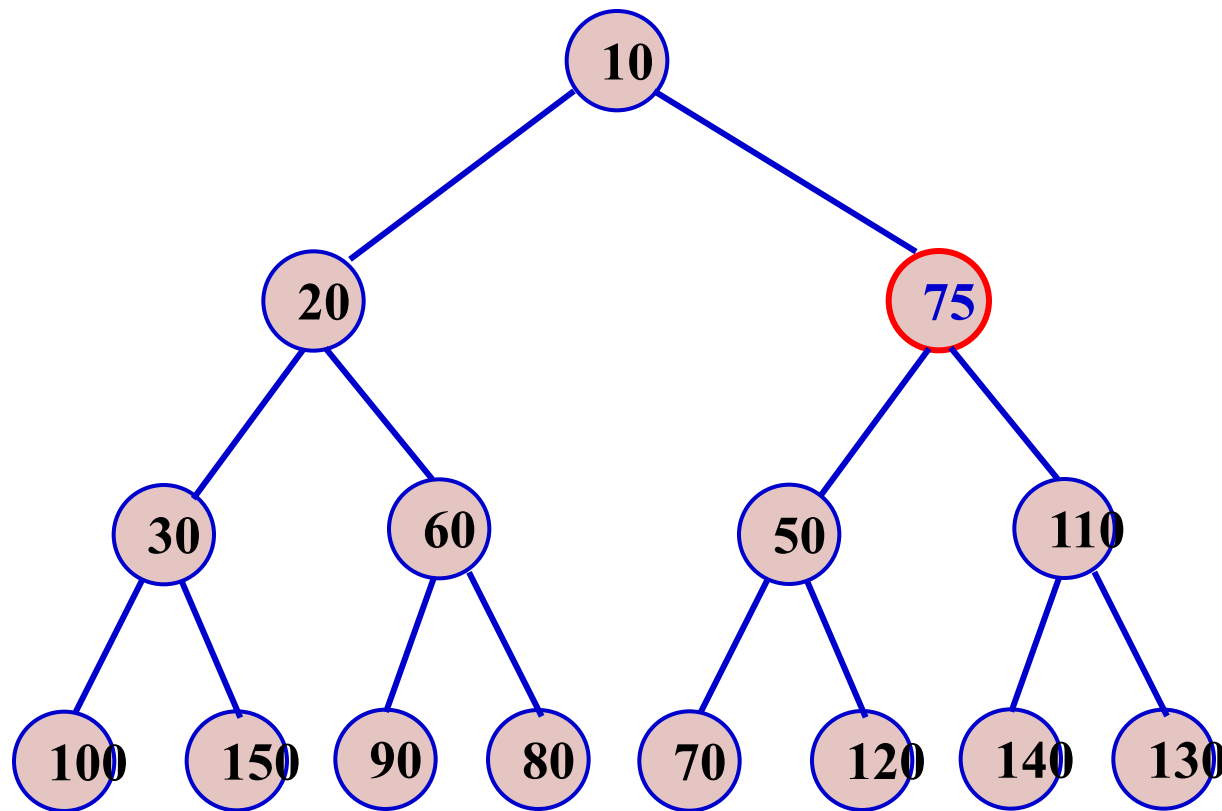
# Heap Operations: Increase Key



***Increase\_key(3, 35, H)***

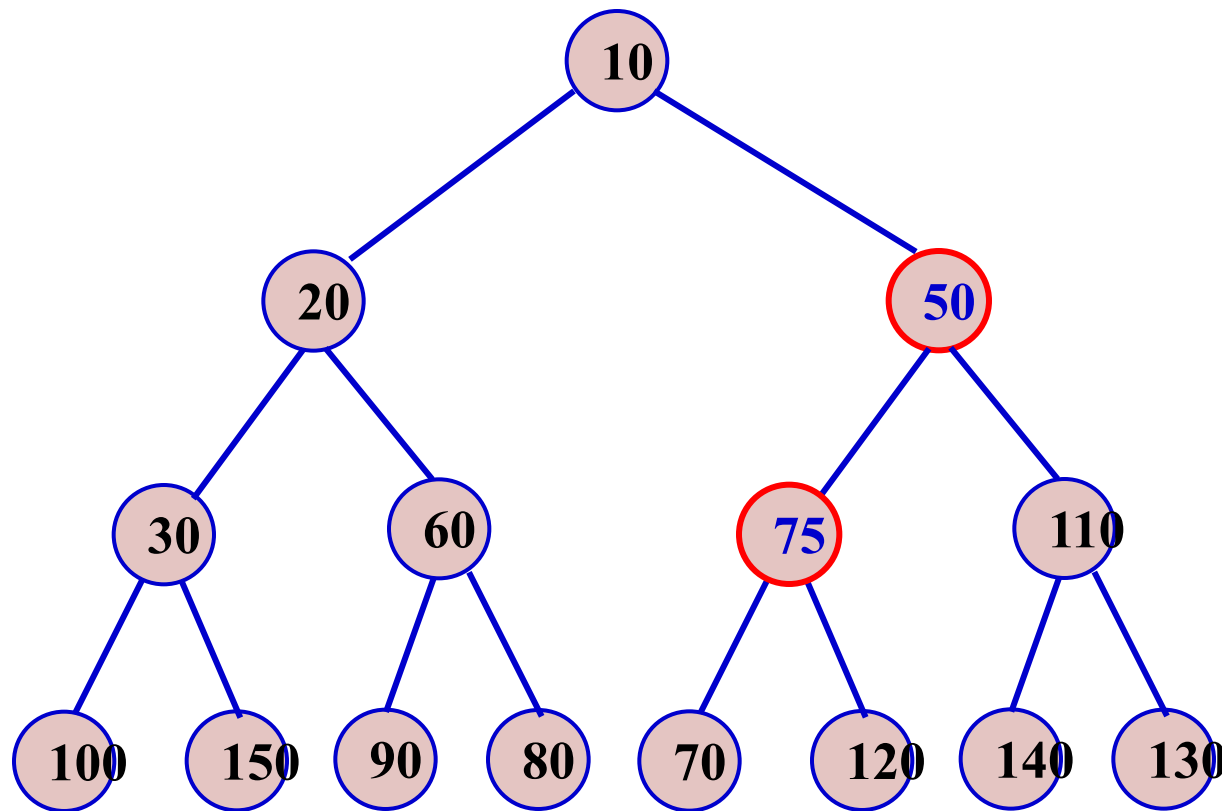


# Heap Operations: Increase Key



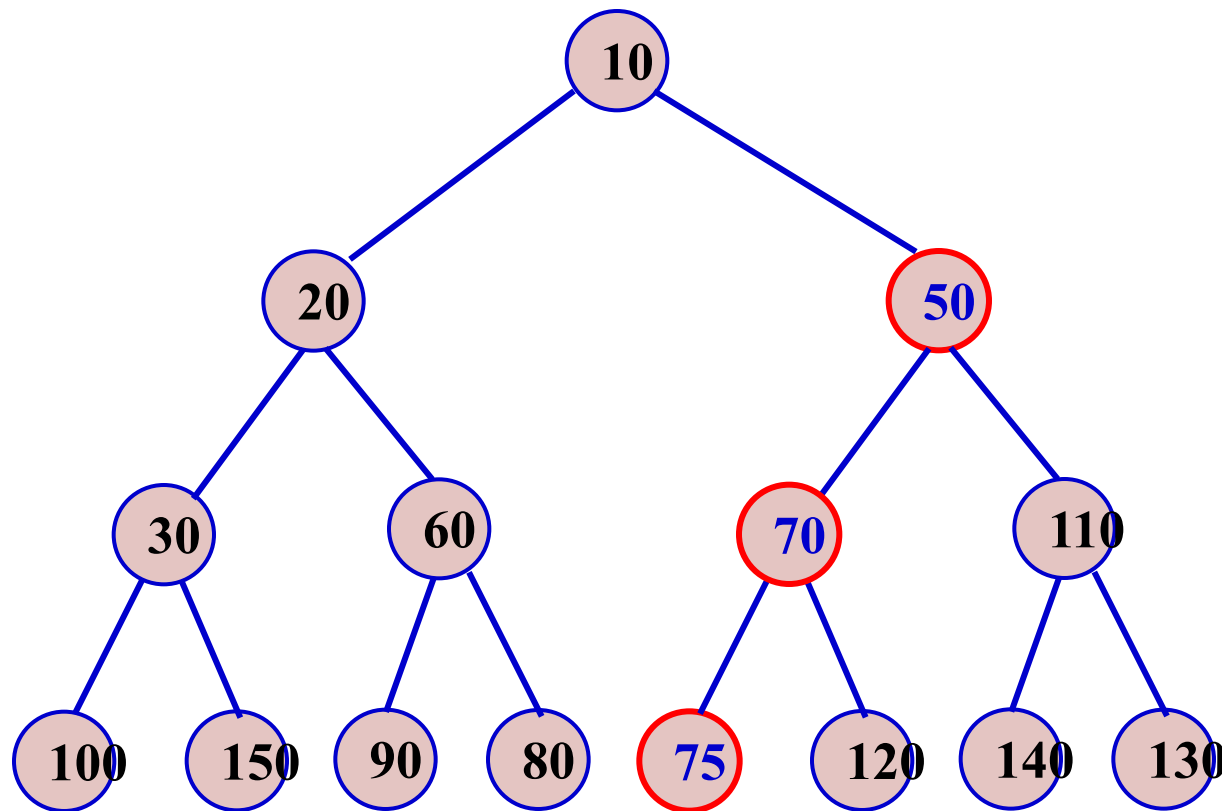
Heap before percolate down(3)

# Heap Operations: Increase Key



**percolate\_down(3)** in action

# Heap Operations: Increase Key

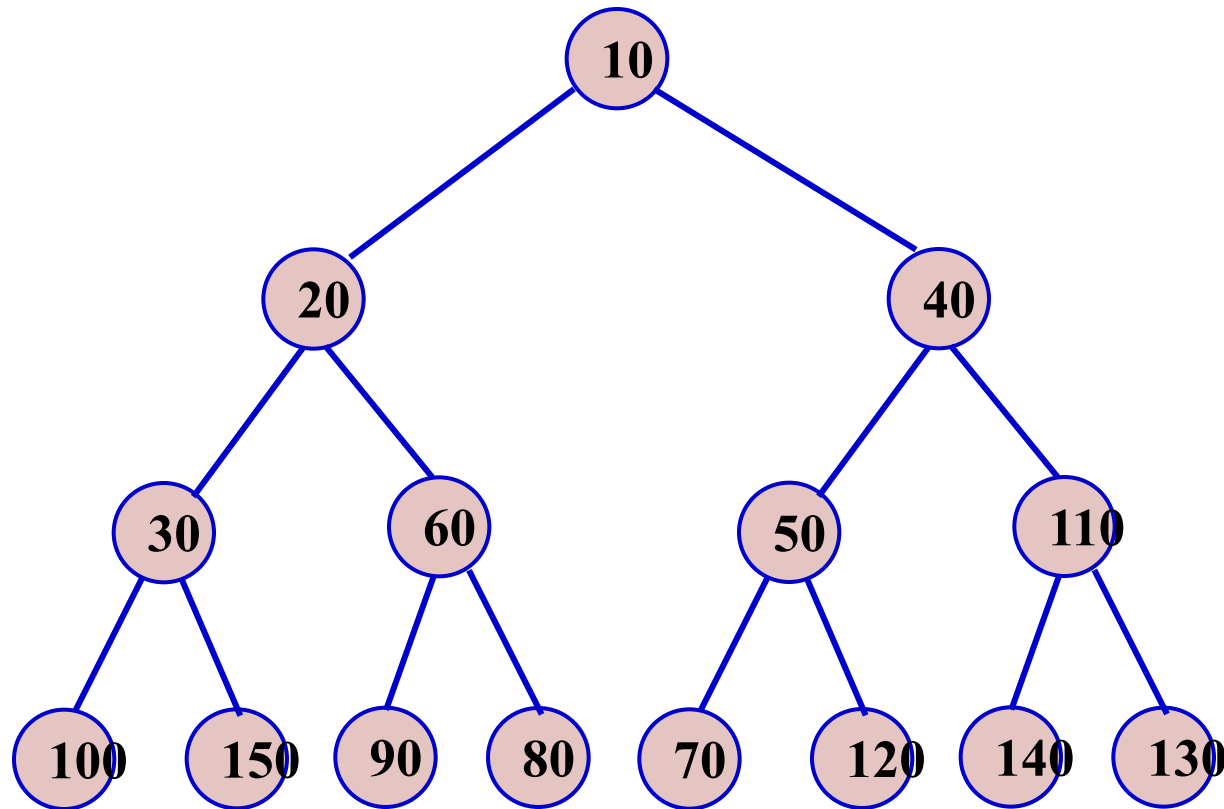


**Heap after percolate\_down(3)**

# Heap Operations: Delete

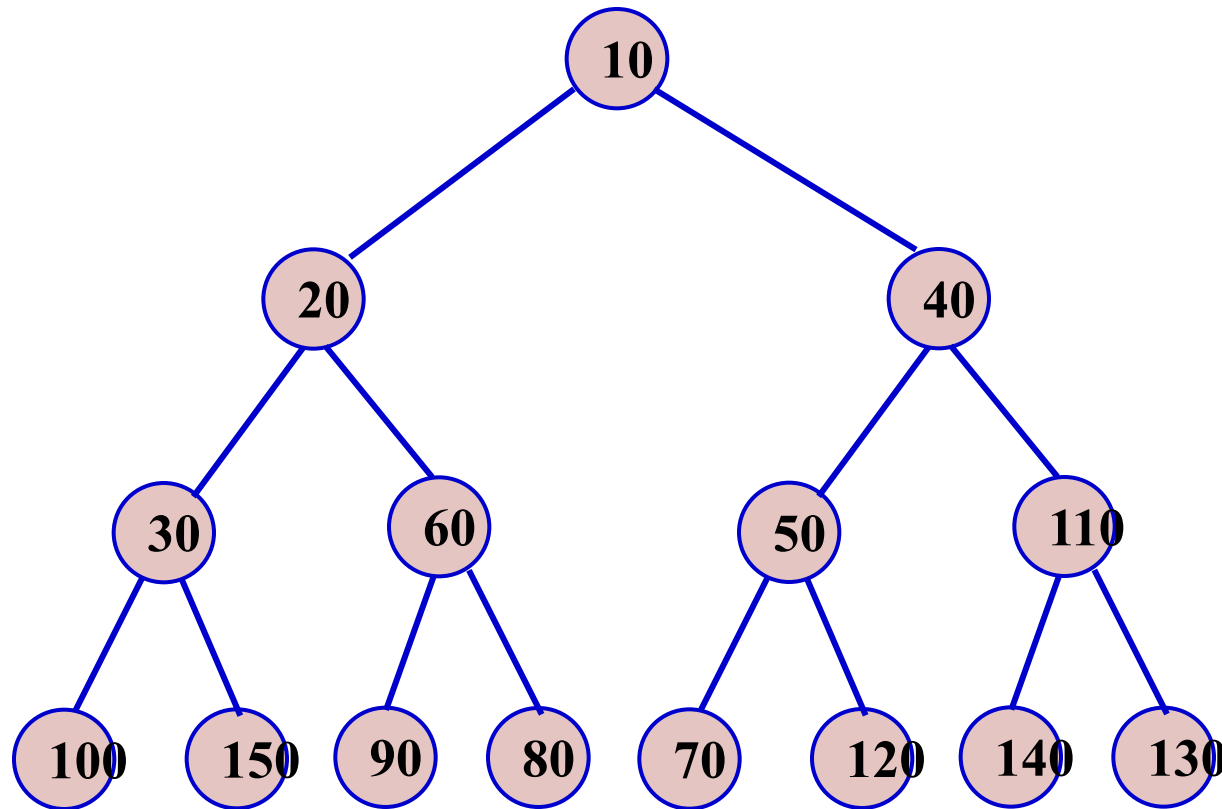
- The ***delete( $x, H$ )*** operation removes the node at position  ***$x$***  from the heap.
- This is done by first performing ***decrease\_key( $x, \infty, H$ )***
- and then performing ***delete\_min( $H$ )***.
- When a process is terminated by a user (instead of finishing normally), it must be removed from the priority queue.

# Heap Operations: Delete



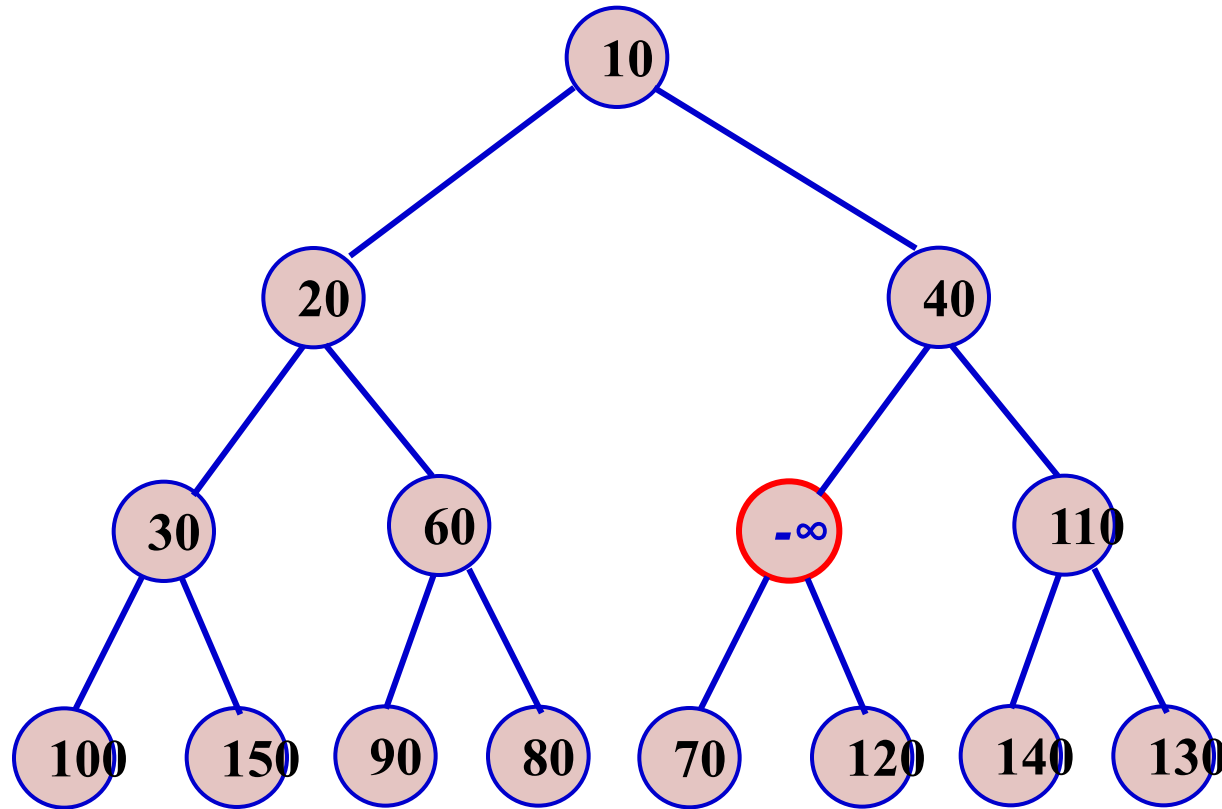
***delete(6, H)***

# Heap Operations: Delete



***delete(6, H) in action: decrease\_key(6,  $\infty$ , H)***

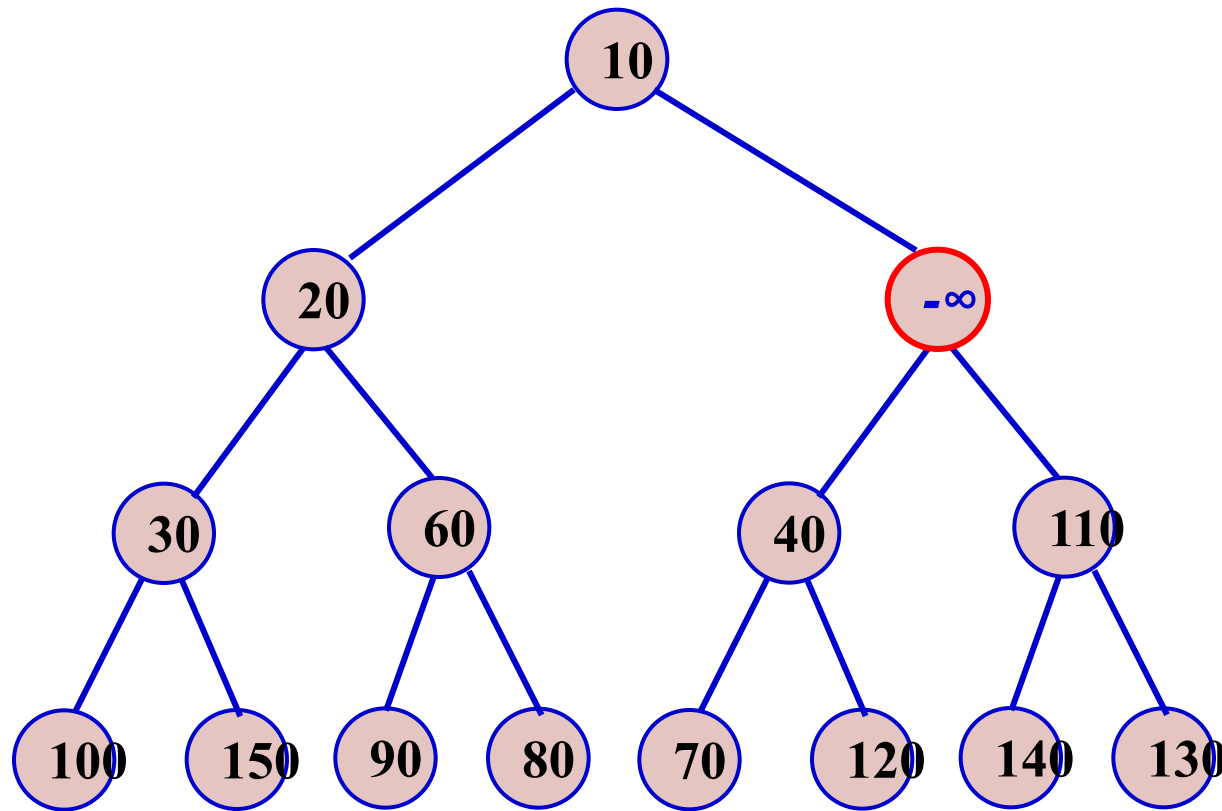
# Heap Operations: Delete



*delete(6, H) in action: decrease\_key(6,  $\infty$ , H)*

**→ percolate\_up(6)**

# Heap Operations: Delete

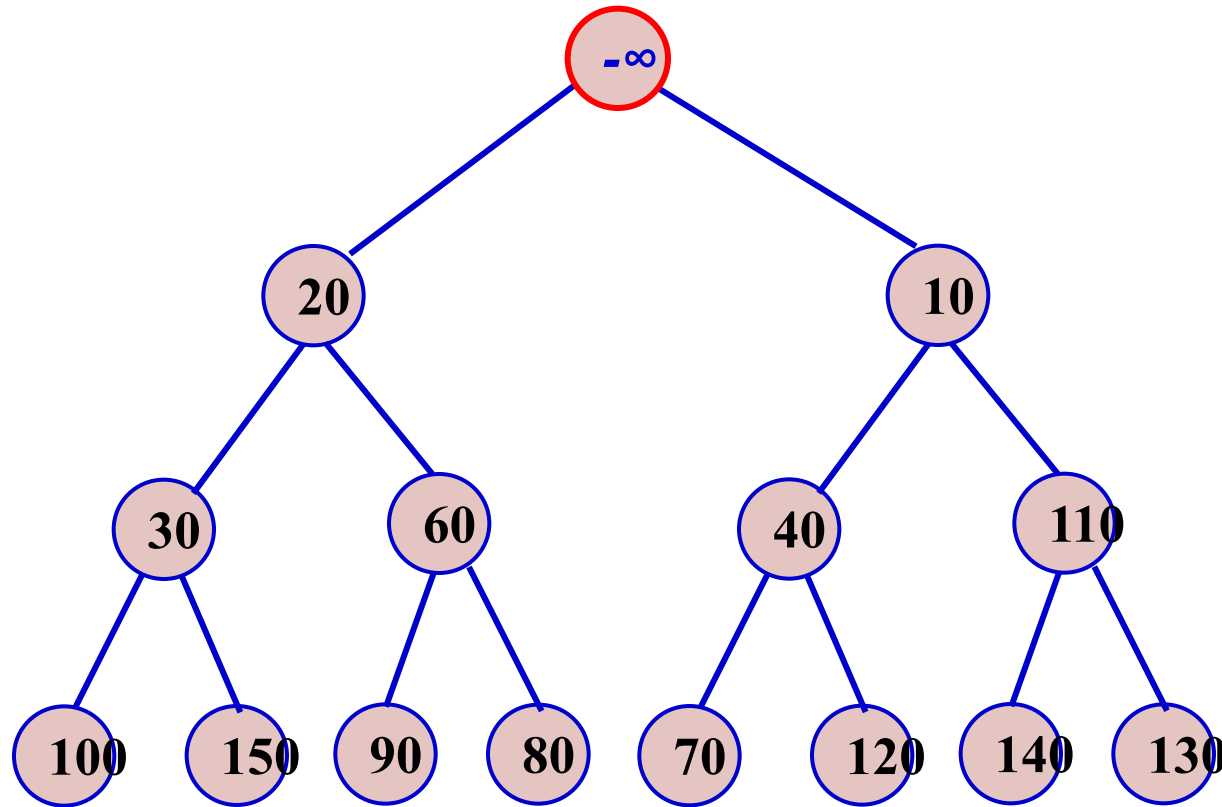


*delete(6, H) in action: decrease\_key(6,  $\infty$ , H)*

**→ percolate\_up(6)**



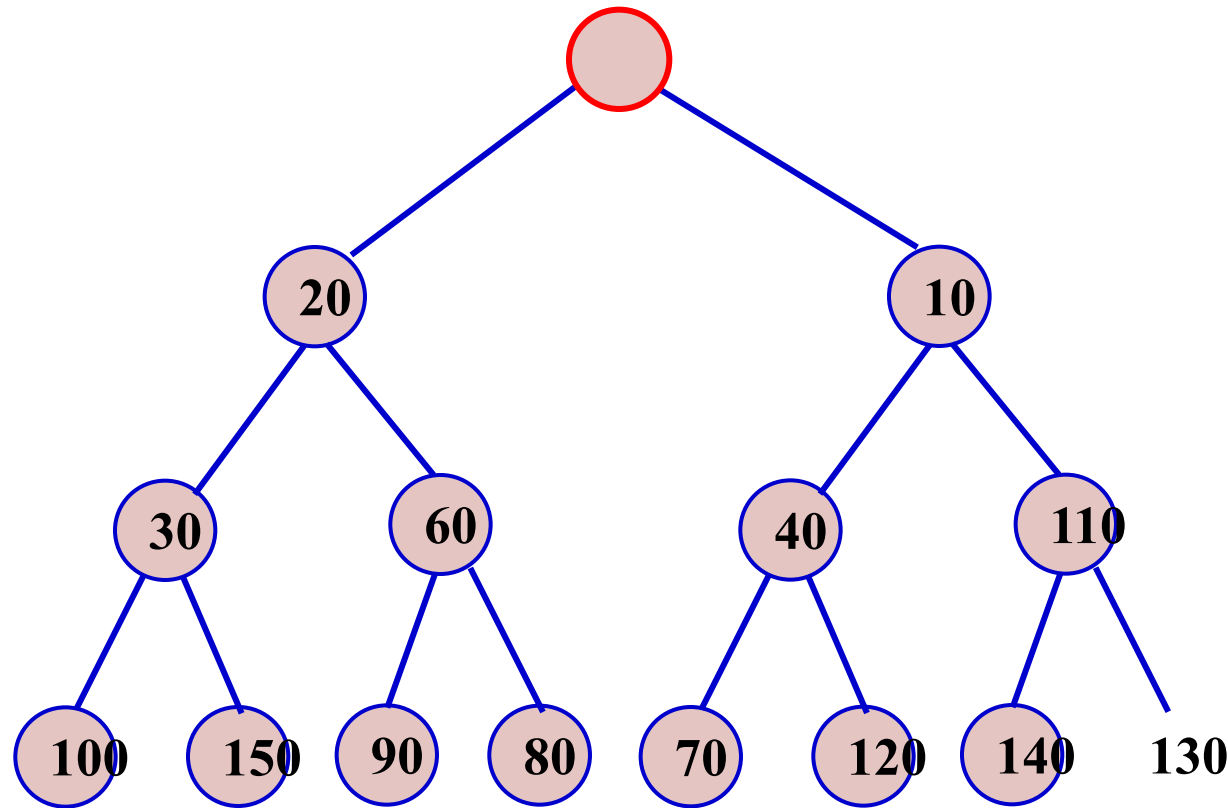
# Heap Operations: Delete



*delete(6, H) in action: decrease\_key(6,  $\infty$ , H)*

**→ delete\_min(H)**

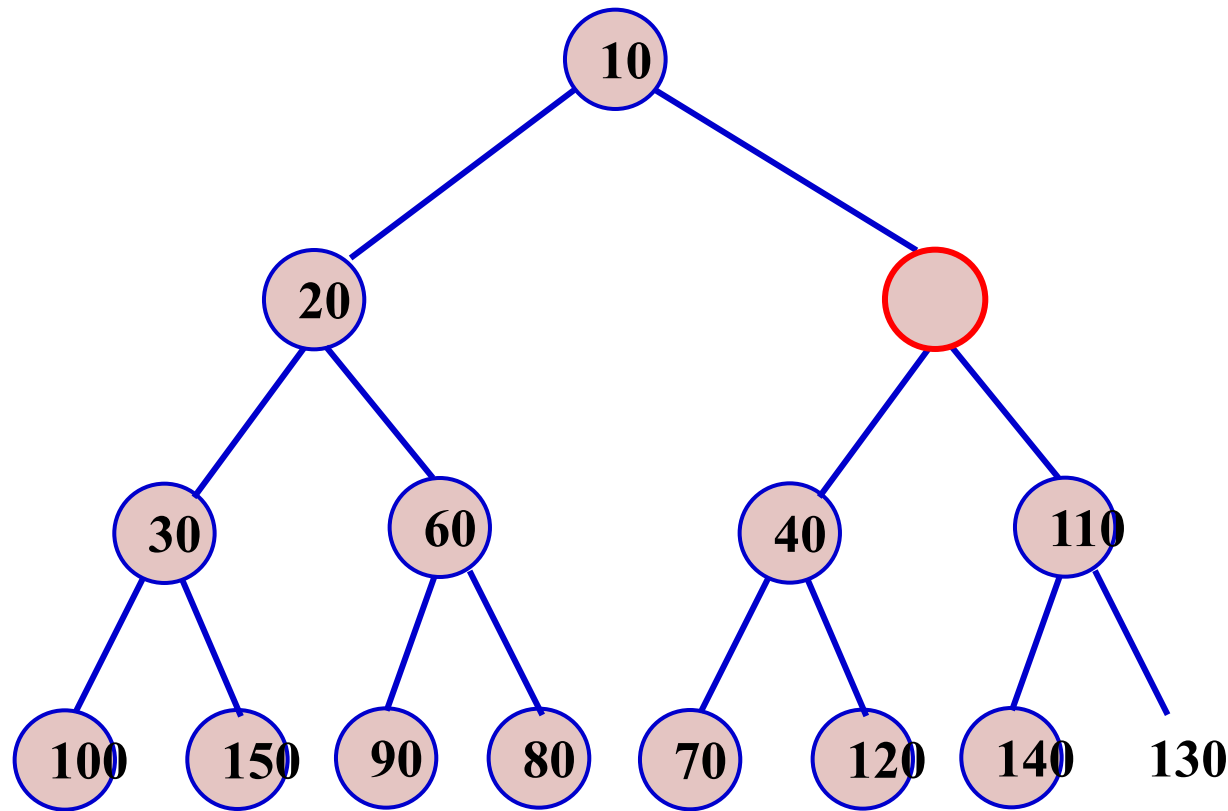
# Heap Operations: Delete



***delete(6, H) in action: decrease\_key(6,  $\infty$ , H)***

**→ delete\_min (H)**

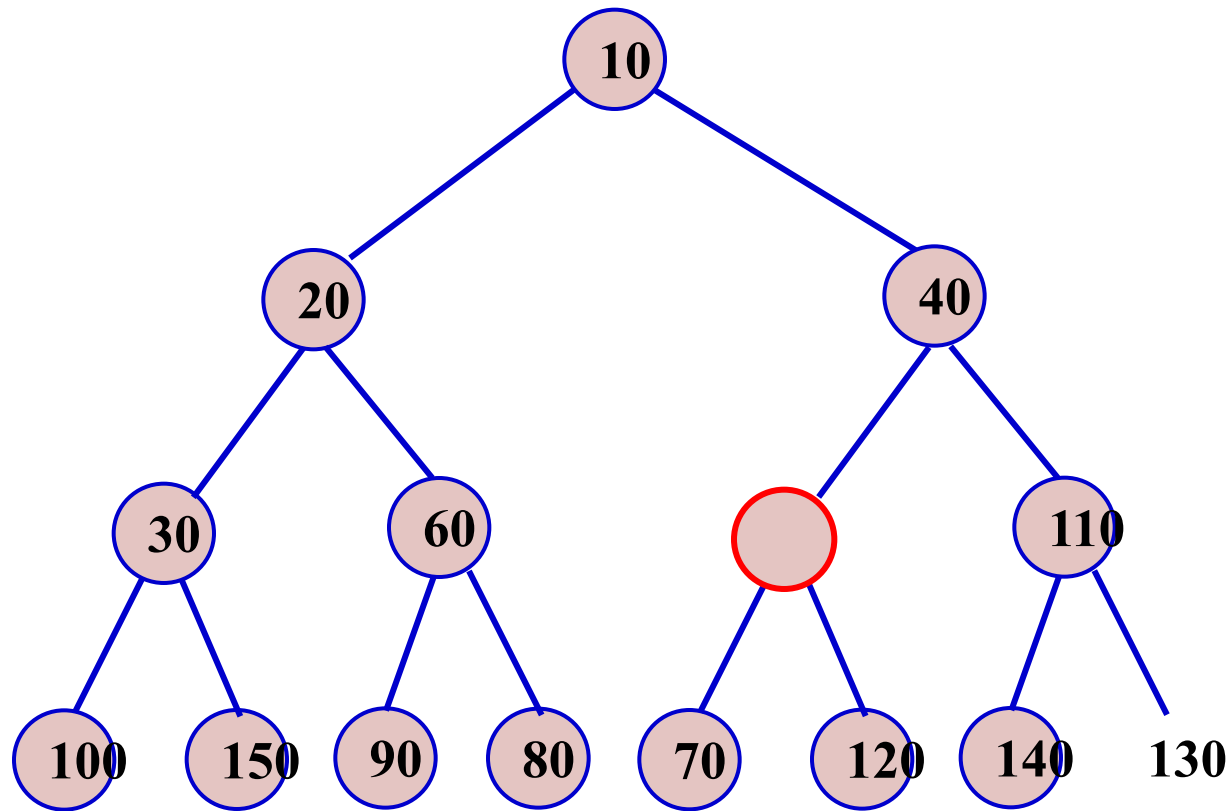
# Heap Operations: Delete



***delete(6, H) in action: decrease\_key(6,  $\infty$ , H)***

**→ delete\_min (H)**

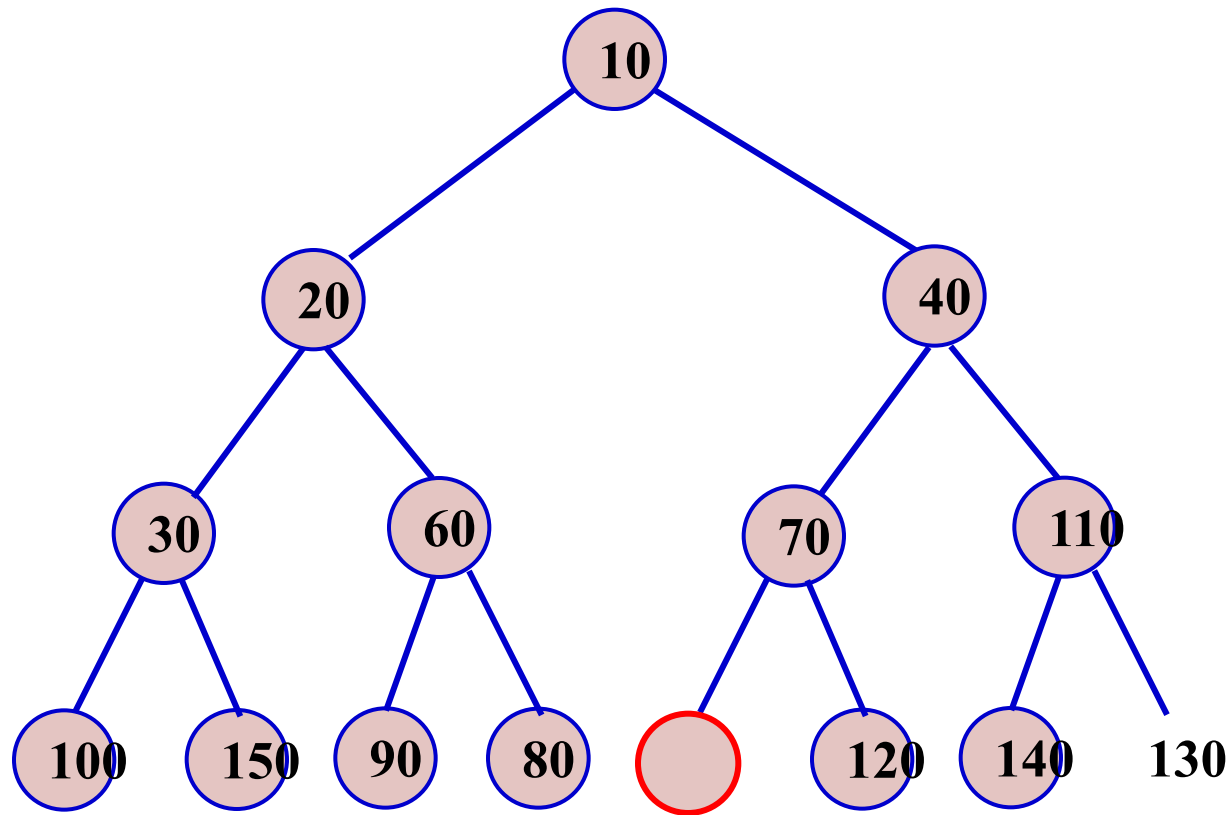
# Heap Operations: Delete



***delete(6, H) in action: decrease\_key(6,  $\infty$ , H)***

***→ delete\_min(H)***

# Heap Operations: Delete



***delete(6, H) in action: decrease\_key(6,  $\infty$ , H)***

**→ delete\_min (H)**

# Questions?

[zahmaad.github.io](https://zahmaad.github.io)